

AWR

AUTOMATED WORKLOAD REPOSITORY

Oracle 10g provides the AWR. The AWR is a repository of performance information collected by the database to aid in the tuning process for DBAs. The main focus for the Oracle database in version 10g is self-manageability. AWR is the place where the data to aid in self-management is stored.

WHAT DOES AWR DO?

Historically people used **bstat** and **estat** to collect Oracle statistics over a time period and then compare them to each other. The bstat/estat approach was replaced with statspack available in Oracle 8i. Statspack was a package provided by Oracle that did roughly the same thing but better. Statspack has now been surpassed in functionality by AWR which is *always* collecting execution statistics for future analysis and tuning performed by all of the expert components provided by Oracle. Oracle recommends that all statspack users switch to AWR in 10g.

The statistics saved in the AWR provide full trending an analysis. In addition, database up/down events will not lose data in the AWR.

AWR collects data in the following categories:

- Base Statistics - general database performance metrics since instance start-up.
- SQL – statistics for each executed SQL statement (# executions, # physical reads, etc)
- Deltas – the rate of change of important stats over time. Similar to our collection technologies that do before and after snapshots and only show the deltas over the specified period of time.
- Expert Advice – results of the expert analysis engine provided in 10g.

The metrics collected by AWR are easily obtained from v\$metricname:

Total Wait Counts	SQL Service Response Time
Total Time Waited	User Limit %
Number of Sessions Waiting (Event)	Current Open Cursors Count
Total Wait Counts	Current Logons Count
Total Time Waited	Global Cache Blocks Lost
Database Time Spent Waiting (%)	Global Cache Blocks Corrupted
Average Users Waiting Counts	Global Cache Average Current Get Time
Host CPU Utilization (%)	Global Cache Average CR Get Time
Database Time Per Sec	GC Current Block Received Per Txn
Txns Per Logon	GC Current Block Received Per Second
Executions Per Sec	GC CR Block Received Per Txn
Executions Per Txn	GC CR Block Received Per Second
Session Limit %	PX downgraded to serial Per Sec
Process Limit %	PX downgraded Parallel Operation Per Sec
PGA Cache Hit %	Sec
Shared Pool Free %	PX downgraded 75% or more Per Sec
Library Cache Miss Ratio	PX downgraded 50% or more Per Sec
Library Cache Hit Ratio	PX downgraded 25% or more Per Sec
Row Cache Miss Ratio	Branch Node Splits Per Txn
Row Cache Hit Ratio	Branch Node Splits Per Sec
Response Time Per Txn	Leaf Node Splits Per Txn
Database CPU Time Ratio	Leaf Node Splits Per Sec
Database Wait Time Ratio	User Rollback Undo Records Applied Per

Txn
 User Rollback UndoRec Applied Per Sec
 CR Undo Records Applied Per Txn
 CR Undo Records Applied Per Sec
 CR Blocks Created Per Txn
 CR Blocks Created Per Sec
 CPU Usage Per Txn
 CPU Usage Per Sec
 Consistent Read Changes Per Txn
 Consistent Read Changes Per Sec
 DB Block Changes Per Txn
 DB Block Changes Per Sec
 Consistent Read Gets Per Txn
 Consistent Read Gets Per Sec
 DB Block Gets Per Txn
 DB Block Gets Per Sec
 Enqueue Requests Per Txn
 Enqueue Requests Per Sec
 Enqueue Deadlocks Per Txn
 Enqueue Deadlocks Per Sec
 Enqueue Waits Per Txn
 Enqueue Waits Per Sec
 Enqueue Timeouts Per Txn
 Enqueue Timeouts Per Sec
 Network Traffic Volume Per Sec
 User Calls Ratio
 Soft Parse Ratio
 Execute Without Parse Ratio
 Rows Per Sort
 Disk Sort Per Txn
 Disk Sort Per Sec
 Cursor Cache Hit Ratio
 Parse Failure Count Per Txn
 Parse Failure Count Per Sec
 Hard Parse Count Per Txn
 Hard Parse Count Per Sec
 Total Parse Count Per Txn
 Total Parse Count Per Sec
 Total Index Scans Per Txn
 Total Index Scans Per Sec
 Full Index Scans Per Txn
 Full Index Scans Per Sec
 Total Table Scans Per Txn
 Total Table Scans Per Sec
 Long Table Scans Per Txn
 Long Table Scans Per Sec
 Redo Writes Per Txn
 Redo Writes Per Sec
 Background Checkpoints Per Sec
 DBWR Checkpoints Per Sec
 Logical Reads Per Txn
 Logical Reads Per Sec
 Recursive Calls Per Txn
 Recursive Calls Per Sec
 User Calls Per Txn
 User Calls Per Sec
 User Rollbacks Percentage
 User Rollbacks Per Sec
 User Commits Percentage
 User Commits Per Sec
 Open Cursors Per Txn

Open Cursors Per Sec
 Logons Per Txn
 Logons Per Sec
 Redo Generated Per Txn
 Redo Generated Per Sec
 Physical Writes Direct Lobs Per Txn
 Physical Writes Direct Lobs Per Sec
 Physical Reads Direct Lobs Per Txn
 Physical Reads Direct Lobs Per Sec
 Physical Writes Direct Per Txn
 Physical Writes Direct Per Sec
 Physical Reads Direct Per Txn
 Physical Reads Direct Per Sec
 Physical Writes Per Txn
 Physical Writes Per Sec
 Physical Reads Per Txn
 Physical Reads Per Sec
 User Transaction Per Sec
 Redo Allocation Hit Ratio
 Memory Sorts Ratio
 Buffer Cache Hit Ratio
 Database Time Per Sec
 Txns Per Logon
 Executions Per Sec
 Executions Per Txn
 Shared Pool Free %
 Library Cache Hit Ratio
 Database CPU Time Ratio
 Consistent Read Changes Per Txn
 Consistent Read Changes Per Sec
 DB Block Changes Per Txn
 DB Block Changes Per Sec
 Consistent Read Gets Per Txn
 Consistent Read Gets Per Sec
 DB Block Gets Per Txn
 DB Block Gets Per Sec
 Soft Parse Ratio
 Execute Without Parse Ratio
 Full Index Scans Per Txn
 Full Index Scans Per Sec
 Total Table Scans Per Txn
 Total Table Scans Per Sec
 Redo Writes Per Txn
 Redo Writes Per Sec
 Logical Reads Per Txn
 Logical Reads Per Sec
 User Calls Per Txn
 User Calls Per Sec
 Logons Per Txn
 Logons Per Sec
 Redo Generated Per Txn
 Redo Generated Per Sec
 Physical Reads Direct Per Txn
 Physical Reads Direct Per Sec
 Physical Writes Per Txn
 Physical Writes Per Sec
 Physical Reads Per Txn
 Physical Reads Per Sec
 User Transaction Per Sec
 Memory Sorts Ratio
 Buffer Cache Hit Ratio

Blocked User Session Count
 Logical Reads Ratio (Sess/Sys) %
 Physical Reads Ratio (Sess/Sys) %
 Total Parse Count (Session)
 Hard Parse Count (Session)
 PGA Memory (Session)
 Physical Reads (Session)
 CPU Time (Session)
 User Transaction Count (Session)

CPU Time Per User Call
 Elapsed Time Per User Call
 Physical Block Writes (Files-Long)
 Physical Block Reads (Files-Long)
 Physical Writes (Files-Long)
 Physical Reads (Files-Long)
 Average File Write Time (Files-Long)
 Average File Read Time (Files-Long)
 Tablespace Space Usage

AWR INSTALLATION

AWR is automatically installed and running with 10g. The new MMON process is responsible for collecting data and populating the AWR.

```
$ ps -ef | grep MID101db
```

oracle	2134	1	0	Feb 16 ?	0:02	ora_cjq0_MID101db
oracle	2136	1	0	Feb 16 ?	0:00	ora_d000_MID101db
oracle	2138	1	0	Feb 16 ?	0:00	ora_s000_MID101db
oracle	4984	1	1	13:51:52 ?	0:00	ora_q000_MID101db
oracle	2150	1	0	Feb 16 ?	0:35	ora_mmon_MID101db
oracle	2152	1	0	Feb 16 ?	0:00	ora_mml1_MID101db
oracle	2154	1	0	Feb 16 ?	1:57	ora_j000_MID101db
oracle	2158	1	0	Feb 16 ?	0:01	ora_j002_MID101db
oracle	2160	1	0	Feb 16 ?	0:00	ora_j003_MID101db

The MMON process takes snapshots of performance data at regular intervals and inserts that data into the AWR tables. The tables containing AWR information are stored in the SYSAUX tablespace (also new in 10G) under the SYS schema. There are a ton of tables in this tablespace – 806 to be exact. However, the AWR related tables all begin with “WR”:

WRM\$_WR_CONTROL	WRI\$_ALERT_THRESHOLD
WRM\$_SNAP_ERROR	WRI\$_ALERT_OUTSTANDING
WRM\$_SNAPSHOT	WRI\$_ALERT_HISTORY
WRM\$_DATABASE_INSTANCE	WRI\$_AGGREGATION_ENABLED
WRM\$_BASELINE	WRI\$_ADV_USAGE
WRI\$_TRACING_ENABLED	WRI\$_ADV_TASKS
WRI\$_SQLSET_STATEMENTS	WRI\$_ADV_SQLW_TABVOL
WRI\$_SQLSET_REFERENCES	WRI\$_ADV_SQLW_TABLES
WRI\$_SQLSET_DEFINITIONS	WRI\$_ADV_SQLW_SUM
WRI\$_SQLSET_BINDS	WRI\$_ADV_SQLW_STMTS
WRI\$_SCH_VOTES	WRI\$_ADV_SQLW_COLVOL
WRI\$_SCH_CONTROL	WRI\$_ADV_SQLT_STATISTICS
WRI\$_OPTSTAT_TAB_HISTORY	WRI\$_ADV_SQLT_RTN_PLAN
WRI\$_OPTSTAT_OPR	WRI\$_ADV_SQLT_PLANS
WRI\$_OPTSTAT_IND_HISTORY	WRI\$_ADV_SQLT_BINDS
WRI\$_OPTSTAT_HISTHEAD_HISTORY	WRI\$_ADV_SQLA_TMP
WRI\$_OPTSTAT_HISTGRM_HISTORY	WRI\$_ADV_SQLA_STMTS
WRI\$_OPTSTAT_AUX_HISTORY	WRI\$_ADV_SQLA_MAP
WRI\$_DBU_USAGE_SAMPLE	WRI\$_ADV_SQLA_FAKE_REG
WRI\$_DBU_HWM_METADATA	WRI\$_ADV_REC_ACTIONS
WRI\$_DBU_HIGH_WATER_MARK	WRI\$_ADV_RECOMMENDATIONS
WRI\$_DBU_FEATURE_USAGE	WRI\$_ADV_RATIONALE
WRI\$_DBU_FEATURE_METADATA	WRI\$_ADV_PARAMETERS

WRI\$_ADV_OBJECTS	WRH\$_SEG_STAT_OBJ
WRI\$_ADV_MESSAGE_GROUPS	WRH\$_SEG_STAT_BL
WRI\$_ADV_JOURNAL	WRH\$_ROWCACHE_SUMMARY_BL
WRI\$_ADV_FINDINGS	WRH\$_RESOURCE_LIMIT
WRI\$_ADV_DIRECTIVES	WRH\$_PGA_TARGET_ADVICE
WRI\$_ADV_DEF_PARAMETERS	WRH\$_PGASTAT
WRI\$_ADV_DEFINITIONS	WRH\$_PARAMETER_NAME
WRI\$_ADV_ACTIONS	WRH\$_PARAMETER_BL
WRH\$_WAITSTAT_BL	WRH\$_OSSTAT_NAME
WRH\$_WAITCLASSMETRIC_HISTORY	WRH\$_OSSTAT_BL
WRH\$_UNDOSTAT	WRH\$_OPTIMIZER_ENV
WRH\$_THREAD	WRH\$_MTTR_TARGET_ADVICE
WRH\$_TEMPSTATXS	WRH\$_METRIC_NAME
WRH\$_TEMPFILE	WRH\$_LOG
WRH\$_TABLESPACE_STAT_BL	WRH\$_LIBRARYCACHE
WRH\$_TABLESPACE_SPACE_USAGE	WRH\$_LATCH_PARENT_BL
WRH\$_SYS_TIME_MODEL_BL	WRH\$_LATCH_NAME
WRH\$_SYSTEM_EVENT_BL	WRH\$_LATCH_MISSES_SUMMARY_BL
WRH\$_SYSSTAT_BL	WRH\$_LATCH_CHILDREN_BL
WRH\$_SYSMETRIC_SUMMARY	WRH\$_LATCH_BL
WRH\$_SYSMETRIC_HISTORY	WRH\$_JAVA_POOL_ADVICE
WRH\$_STAT_NAME	WRH\$_INSTANCE_RECOVERY
WRH\$_SQL_WORKAREA_HISTOGRAM	WRH\$_FILESTATXS_BL
WRH\$_SQL_SUMMARY	WRH\$_FILEMETRIC_HISTORY
WRH\$_SQL_PLAN	WRH\$_EVENT_NAME
WRH\$_SQLTEXT	WRH\$_ENQUEUE_STAT
WRH\$_SQLSTAT_BL	WRH\$_DLM_MISC_BL
WRH\$_SQLBIND_BL	WRH\$_DB_CACHE_ADVICE_BL
WRH\$_SHARED_POOL_ADVICE	WRH\$_DATAFILE
WRH\$_SGASTAT_BL	WRH\$_CURRENT_BLOCK_SERVER
WRH\$_SGA	WRH\$_CR_BLOCK_SERVER
WRH\$_SESSMETRIC_HISTORY	WRH\$_CLASS_CACHE_TRANSFER_BL
WRH\$_SERVICE_WAIT_CLASS_BL	WRH\$_BUFFER_POOL_STATISTICS
WRH\$_SERVICE_STAT_BL	WRH\$_BG_EVENT_SUMMARY
WRH\$_SERVICE_NAME	WRH\$_ACTIVE_SESSION_HISTORY_BL

The third letter of each table name signifies the type of data that it contains.

- I – advisory functions (SQL Advice, Space Advice, etc)
- M – metadata information
- H – historical data

Historical data is already populated in the tables. For instance, a query against WRH\$_TABLESPACE_SPACE_USAGE yields 1032 rows even though only 6 tablespaces exist in the database. This is due to data being captured and logged at regular intervals. Any tool would have everything it needs to chart tablespace space usage over time.

Oracle also adds views on top of these base tables. The views all begin with DBA_HIST:

DBA_HIST_DATABASE_INSTANCE	DBA_HIST_TEMPFILE
DBA_HIST_SNAPSHOT	DBA_HIST_TEMPSTATXS
DBA_HIST_SNAP_ERROR	DBA_HIST_SQLSTAT
DBA_HIST_BASELINE	DBA_HIST_SQLTEXT
DBA_HIST_WR_CONTROL	DBA_HIST_SQL_SUMMARY
DBA_HIST_DATAFILE	DBA_HIST_SQL_PLAN
DBA_HIST_FILESTATXS	DBA_HIST_SQLBIND

DBA_HIST_OPTIMIZER_ENV	DBA_HIST_SYS_TIME_MODEL
DBA_HIST_EVENT_NAME	DBA_HIST_OSSTAT_NAME
DBA_HIST_SYSTEM_EVENT	DBA_HIST_OSSTAT
DBA_HIST_BG_EVENT_SUMMARY	DBA_HIST_PARAMETER_NAME
DBA_HIST_WAITSTAT	DBA_HIST_PARAMETER
DBA_HIST_ENQUEUE_STAT	DBA_HIST_UNDOSTAT
DBA_HIST_LATCH_NAME	DBA_HIST_SEG_STAT
DBA_HIST_LATCH	DBA_HIST_SEG_STAT_OBJ
DBA_HIST_LATCH_CHILDREN	DBA_HIST_METRIC_NAME
DBA_HIST_LATCH_PARENT	DBA_HIST_SYSMETRIC_HISTORY
DBA_HIST_LATCH_MISSES_SUMMARY	DBA_HIST_SYSMETRIC_SUMMARY
DBA_HIST_LIBRARYCACHE	DBA_HIST_SESSMETRIC_HISTORY
DBA_HIST_DB_CACHE_ADVICE	DBA_HIST_FILEMETRIC_HISTORY
DBA_HIST_BUFFER_POOL_STAT	DBA_HIST_WAITCLASSMET_HISTORY
DBA_HIST_ROWCACHE_SUMMARY	DBA_HIST_DLM_MISC
DBA_HIST_SGA	DBA_HIST_CR_BLOCK_SERVER
DBA_HIST_SGASTAT	DBA_HIST_CURRENT_BLOCK_SERVER
DBA_HIST_PGASTAT	DBA_HIST_CLASS_CACHE_TRANSFER
DBA_HIST_RESOURCE_LIMIT	DBA_HIST_ACTIVE_SESS_HISTORY
DBA_HIST_SHARED_POOL_ADVICE	DBA_HIST_TABLESPACE_STAT
DBA_HIST_SQL_WORKAREA_HSTGRM	DBA_HIST_LOG
DBA_HIST_PGA_TARGET_ADVICE	DBA_HIST_MTTR_TARGET_ADVICE
DBA_HIST_INSTANCE_RECOVERY	DBA_HIST_TBSPC_SPACE_USAGE
DBA_HIST_JAVA_POOL_ADVICE	DBA_HIST_SERVICE_NAME
DBA_HIST_THREAD	DBA_HIST_SERVICE_STAT
DBA_HIST_STAT_NAME	DBA_HIST_SERVICE_WAIT_CLASS
DBA_HIST_SYSSTAT	

The frequency of data collection is 30 minutes by default but that can be adjusted. All functionality for driving the workload repository is done via the Oracle supplied package DBMS_WORKLOAD_REPOSITORY. The package header spec is below:

```
-- ***** --
-- DBMS_WORKLOAD_REPOSITORY Constants
-- ***** --

-- Minimum and Maximum values for the
-- Snapshot Interval Setting (in minutes)
MIN_INTERVAL      CONSTANT NUMBER := 10;           /* 10 minutes */
MAX_INTERVAL      CONSTANT NUMBER := 52560000;     /* 100 years */

-- Minimum and Maximum values for the
-- Snapshot Retention Setting (in minutes)
MIN_RETENTION     CONSTANT NUMBER := 1440;         /* 1 day */
MAX_RETENTION     CONSTANT NUMBER := 52560000;     /* 100 years */

-- ***** --
-- DBMS_WORKLOAD_REPOSITORY Routines
-- ***** --

--
-- create snapshot()
--   Creates a snapshot in the workload repository.
--
--   This routine will come in two forms: procedure and function.
--   The function returns the snap id for the snapshot just taken.
--
-- Input arguments:
--   flush level          - flush level for the snapshot:
--                         either 'TYPICAL' or 'ALL'
--
-- Returns:
--   NUMBER               - snap id for snapshot just taken.
--
```

```

PROCEDURE create_snapshot(flush_level IN VARCHAR2 DEFAULT 'TYPICAL'
                        );

FUNCTION create_snapshot(flush_level IN VARCHAR2 DEFAULT 'TYPICAL'
                        ) RETURN NUMBER;

--
-- drop_snapshot_range()
-- purge the snapshots for the given range of snapshots.
--
-- Input arguments:
--   low snap id      - low snapshot id of snapshots to drop
--   high snap id     - high snapshot id of snapshots to drop
--   dbid             - database id (default to local DBID)
--
PROCEDURE drop_snapshot_range(low_snap_id      IN NUMBER,
                             high_snap_id     IN NUMBER,
                             dbid             IN NUMBER DEFAULT NULL
                             );

--
-- modify_snapshot_settings()
-- Procedure to adjust the settings of the snapshot collection.
--
-- Input arguments:
--   retention        - new retention time (in minutes). The
--                     specified value must be in the range:
--                     MIN_RETENTION (1 day) to
--                     MAX_RETENTION (100 years)
--
--                     If ZERO is specified, snapshots will be
--                     retained forever. A large system-defined
--                     value will be used as the retention setting.
--
--                     If NULL is specified, the old value for
--                     retention is preserved.
--
--   interval        - the interval between each snapshot, in
--                     units of minutes. The specified value
--                     must be in the range:
--                     MIN_INTERVAL (10 minutes) to
--                     MAX_INTERVAL (100 years)
--
--                     If ZERO is specified, automatic and manual
--                     snapshots will be disabled. A large
--                     system-defined value will be used as the
--                     interval setting.
--
--                     If NULL is specified, the
--                     current value is preserved.
--
--   dbid            - database identifier for the database to
--                     adjust setting. If NULL is specified, the
--                     local dbid will be used.
--
-- For example, the following statement can be used to set the
-- Retention and Interval to their minimum settings:
--
--   dbms_workload_repository.modify_snapshot_settings
--   (retention => DBMS_WORKLOAD_REPOSITORY.MIN_RETENTION
--    interval  => DBMS_WORKLOAD_REPOSITORY.MIN_INTERVAL)
--
-- The following statement can be used to set the Retention to
-- 7 days and the Interval to 60 minutes:
--
--   dbms_workload_repository.modify_snapshot_settings
--   (retention => 10080, interval => 60);
--
PROCEDURE modify_snapshot_settings(retention IN NUMBER DEFAULT NULL,
                                  interval  IN NUMBER DEFAULT NULL,
                                  dbid      IN NUMBER DEFAULT NULL
                                  );

```

```

--
-- create baseline()
-- Routine to create a baseline. A baseline is set of
-- of statistics defined by a (begin, end) pair of snapshots.
--
-- This routine will come in two forms: procedure and function.
-- The function returns the baseline_id for the baseline just created.
--
-- Input arguments:
-- start_snap_id      - start snapshot sequence number for baseline
-- end_snap_id        - end snapshot sequence number for baseline
-- baseline_name      - name of baseline (required)
-- dbid               - optional dbid, default to Local DBID
--
-- Returns:
-- NUMBER              - baseline_id for the baseline just created
--

PROCEDURE create_baseline(start_snap_id IN NUMBER,
                        end_snap_id   IN NUMBER,
                        baseline_name  IN VARCHAR2,
                        dbid           IN NUMBER DEFAULT NULL
                        );

FUNCTION create_baseline(start_snap_id IN NUMBER,
                        end_snap_id   IN NUMBER,
                        baseline_name  IN VARCHAR2,
                        dbid           IN NUMBER DEFAULT NULL
                        ) RETURN NUMBER;

--
-- drop_baseline()
-- drops a baseline (by name)
--
-- Input arguments:
-- baseline_name      - name of baseline to drop
-- dbid               - database id, default to local DBID
-- cascade            - if TRUE, the range of snapshots associated
--                     with the baseline will also be dropped.
--                     Otherwise, only the baseline is removed.
--
PROCEDURE drop_baseline(baseline_name IN VARCHAR2,
                        cascade        IN BOOLEAN DEFAULT false,
                        dbid           IN NUMBER DEFAULT NULL
                        );

-- *****
-- awr_report_text and_html (FUNCTION)
-- This is the table function that will display the
-- AWR report in either text or HTML. The output will be
-- one column of VARCHAR2(80) or (150), respectively
--
-- The report will take as input the following parameters:
-- l_dbid      - database identifier
-- l_inst_num  - instance number
-- l_bid       - Begin Snap Id
-- l_eid       - End Snapshot Id
-- *****
FUNCTION awr_report_text(l_dbid      IN NUMBER,
                        l_inst_num  IN NUMBER,
                        l_bid       IN NUMBER,
                        l_eid       IN NUMBER,
                        l_options   IN NUMBER DEFAULT 0)
RETURN awrrpt text type table PIPELINED;

FUNCTION awr_report_html(l_dbid      IN NUMBER,
                        l_inst_num  IN NUMBER,
                        l_bid       IN NUMBER,
                        l_eid       IN NUMBER,
                        l_options   IN NUMBER DEFAULT 0)
RETURN awrrpt html type table PIPELINED;

END dbms_workload_repository;

```

Manual snapshots are also available by using the *create_snapshot* function within the package. The execution of this package produces the following TKPROF results:

OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	127	0.10	0.10	0	0	0	0
Execute	214	1.93	2.40	43	1510	7022	2213
Fetch	229	0.03	0.04	2	471	0	212
total	570	2.06	2.55	45	1981	7022	2425

Now we can see that the total number of statements on an AWR snapshot refresh:

```
1 session in tracefile.
7 user SQL statements in trace file.
127 internal SQL statements in trace file.
134 SQL statements in trace file.
113 unique SQL statements in trace file.
4057 lines in trace file.
31 elapsed seconds in trace file.
```

134 SQL statements submitted. Some statements even use the RULE hint - go figure. This is the exact same architecture of DBXray - a package that executes at regular intervals to insert data into tables. I'm not sure why MMON process is needed unless it only does the scheduling.

LIMITATIONS

It can only be assumed that this data capture and storage utility provided by the MMON process utilizes CPU. However, during my tests of manually creating snapshots, I could never get MMON to use more the .5% of the CPU although this is a test database.