





Using AWR For Memory Analysis

Mike Ault, Oracle Guru
May, 2011



Michael R. Ault, Oracle Guru

- Nuclear Navy 6 years
- Nuclear Chemist/Programmer 10 years
- Kennedy Western University Graduate
- Bachelors Degree Computer Science
- Certified in all Oracle Versions Since 6
- Oracle DBA, author, since 1990








ORACLE

**CERTIFIED
PROFESSIONAL**



Books by Michael R. Ault



StatspackAnalyzer.com

Free Statspack/AWR Analysis

Sponsored by Texas Memory Systems

- Looks for IO bottlenecks and other configuration issues.
- Straightforward tuning advice





Preparation for Analysis

- Know your systems normal performance fingerprint
- Be familiar with Concepts and Tuning Guides
- Have “normal” AWR/Statspacks for comparison



Oracle and memory

- DB cache – other than direct read/write, all data, index and undo go through here
- Shared Pool – SQL area, PL/SQL library, dictionary caches and lots more
- Streams pool – Only if streams are used
- Java Pool – usually small
- PGA – Each process gets a PGA, stores cursor and other process related information
- Log buffers – Circular buffers for redo information



DB Cache

- Default – should be largest area
- Recycle – for frequently scanned large objects
- Keep – For frequently accessed small objects
- 2-32K areas – Was originally for TTS, now used to tune items (usually in RAC)



Shared Pool

- Shared SQL, PL/SQL, dictionary cache plus
- 34 other areas in 9i
- 551 in 10gR2 (non-RAC) 670 (with RAC)
- 878 in 11gR2 (non-RAC)
 - Report only shows those that change
 - There have been bugs with space leaks in shared pool sub-pools



Top-Down Approach

- Report starts with settings overview
- Next provides Top-5 waits
- Use the Waits to guide further investigation



AWR Report Header

WORKLOAD REPOSITORY report for

DB Name	DB Id	Instance	Inst Num	Startup Time	Release	RAC
AULTDB	4030696936	aultdb1	1	04-Aug-08 10:16	11.1.0.6.0	YES
Host Name	Platform			CPU	Cores	Sockets
Memory(GB)						
-----	-----	-----	-----	-----	-----	-----
aultlinux3	Linux IA (32-bit)		2	1	1	2.97
Snap Id	Snap Time	Sessions	Curs/Sess			
-----	-----	-----	-----	-----	-----	-----
Begin Snap:	91 04-Aug-08 12:00:15	41	1.2			
End Snap:	92 04-Aug-08 13:00:28	47	1.1			
Elapsed:	60.22 (mins)					
DB Time:	139.52 (mins)					
Cache Sizes	Begin	End				
-----	-----	-----				
Buffer Cache:	1,312M	1,312M	Std Block Size:		8K	
Shared Pool Size:	224M	224M	Log Buffer:		10,604K	



Signs of Memory Issues

- High sequential reads
- Excessive library latches
- Large number of sorts/hashes/GTT/bitmap ops to disk
- Large amount of IO to the temporary tablespace
- Buffer busy waits with free buffer waits
- Indications in Cache, shared, streams or java pool advisors
- Excessive reparsing and reloads of SQL and PL/SQL
- High percentage of use in the shared pool
- High CPU cycles



Load Profile Section

Load Profile ~~~~~	Per Second -----	Per Transaction -----	Per Exec -----	Per Call -----
DB Time(s):	2.3	7.1	0.63	1.05
DB CPU(s):	0.3	0.9	0.07	0.13
Redo size:	800.5	2,461.8		
Logical reads:	6,307.6	19,396.7		
Block changes:	3.6	10.9		
Physical reads:	2,704.9	8,317.8		
Physical writes:	86.9	267.3		
User calls:	2.2	6.8		
Parses:	2.0	6.1		
Hard parses:	0.0	0.1		
W/A MB processed:	932,965.4	2,868,990.9		
Logons:	0.1	0.2		
Executes:	3.7	11.3		
Rollbacks:	0.1	0.3		
Transactions:	0.3			



What Are Your Efficiencies

- Should be close to 100%
- Parse issues usually are a result of:
 - Bad bind variable usage
 - Insufficient memory
 - Will also be co-indicated by low percentage of memory for multiple SQL execution



Load Profile Section

Instance Efficiency Percentages (Target 100%)

```
~~~~~
      Buffer Nowait %: 100.00      Redo NoWait %: 99.97
      Buffer Hit %: 96.09      In-memory Sort %: 100.00
      Library Hit %: 98.17      Soft Parse %: 97.88
      Execute to Parse %: 45.80      Latch Hit %: 99.95
Parse CPU to Parse Elapsed %: 0.00      % Non-Parse CPU: 99.77
```

```
Shared Pool Statistics      Begin      End
      -----      -----
      Memory Usage %: 81.53      85.39
      % SQL with executions>1: 79.29      79.48
      % Memory for SQL w/exec>1: 76.73      78.19
```



Top 5 Waits Section

- Critical to look closely at this section
- Use highest wait times to guide investigation
 - DB FILE type waits – physical IO
 - BUFFER type waits – Logical IO
 - LOG type waits – Redo related
 - PX – Parallel Query
 - GC – Global Cache (RAC related)
 - Undo – Undo or rollback segment related



Top 5 waits section With possible cache starvation

Top 5 Timed Foreground Events

~~~~~

| Event                   | Waits   | Time(s) | Avg<br>wait<br>(ms) | % DB<br>time | Wait Class |
|-------------------------|---------|---------|---------------------|--------------|------------|
| db file sequential read | 465,020 | 3,969   | 9                   | 47.4         | User I/O   |
| DB CPU                  |         | 995     |                     | 11.9         |            |
| db file parallel read   | 2,251   | 322     | 143                 | 3.8          | User I/O   |
| db file scattered read  | 15,268  | 153     | 10                  | 1.8          | User I/O   |
| gc current block 2-way  | 108,739 | 116     | 1                   | 1.4          | Cluster    |





# Top 5 Waits Section With Shared Pool Issues

## Top 5 Timed Events

| ~~~~~                    |             |          | Avg  | %Total |      |            |
|--------------------------|-------------|----------|------|--------|------|------------|
|                          |             |          | wait | Call   |      |            |
| Event                    | Waits       | Time (s) | (ms) | Time   | Wait | Class      |
| -----                    |             |          |      |        |      |            |
| CPU time                 |             | 435,461  |      | 41.1   |      |            |
| PX Deq Credit: send blkd | 124,829,330 | 138,223  | 1    | 13.0   |      | Other      |
| library cache pin        | 20,347      | 57,692   | 2835 | 5.4    |      | Concurrenc |
| library cache lock       | 19,226      | 56,078   | 2917 | 5.3    |      | Concurrenc |
| db file sequential read  | 16,798,329  | 42,215   | 3    | 4.0    |      | User I/O   |

## Top 5 Timed Events

| ~~~~~                   |           |          | Avg  | %Total |      |            |
|-------------------------|-----------|----------|------|--------|------|------------|
|                         |           |          | wait | Call   |      |            |
| Event                   | Waits     | Time (s) | (ms) | Time   | Wait | Class      |
| -----                   |           |          |      |        |      |            |
| CPU time                |           | 24,956   |      | 29.3   |      |            |
| latch: library cache    | 1,757,331 | 9,886    | 6    | 11.6   |      | Concurrenc |
| db file sequential read | 759,605   | 6,146    | 8    | 7.2    |      | User I/O   |
| cursor: pin S           | 2,103,389 | 4,988    | 2    | 5.9    |      | Other      |
| log file sync           | 250,039   | 2,387    | 10   | 2.8    |      | Commit     |



# Buffer Type Waits

- latch: cache buffers chains – Hot blocks, check for hot objects, high IO rates
- free buffer waits – Insufficient buffers, processes holding buffers too long, IO subsystem over loaded
- buffer busy waits – See what is causing them further along in report
- gc buffer busy – Overloaded interconnect, find problem objects and tune
- log buffer space – High load, too small a log buffer, increase log buffer size
- latch: cache buffers lru chain – Freelist issues, hot blocks, new buffers, buffers being written
- latch: cache buffer handles – Freelist issues, hot blocks
- buffer busy - See what is causing them further along in report
- no free buffers – Insufficient buffers, dbwr contention
- Free buffer waits – insufficient buffers



# Fixing Cache Waits

- Reduce logical IO rates (buffer caches latch)
- Increase the cache size (lru chain latch)
- Increase the cache size (free buffer waits)
- Increase `_db_block_lru_latches`
- Increase `_db_block_hash_buckets`
- Reduce hot blocks



# Shared Pool Waits

- library cache pin – Loading or compiling same SQL
  - library cache lock – Loading or compiling same SQL
  - \*latch: library cache – Usually a result of excessive parsing
  - latch: shared pool latch – Parsing issues
  - \*latch: library cache lock – Usually a result of excessive parsing
  - \*latch: library cache – Usually a result of excessive parsing
  - row cache lock – shared pool too small
  - Library cache load lock – Wait for a reload by another session. Excessive hard/soft parsing.
- \* Gone in 11g to mutex



# Shared Pool Mutexes

- Cursor:mutex X – resource is busy, requestor needs exclusive access
- Cursor:mutex S – resource is held in X mode by another session
- Cursor:pin X – resource is held in S or X by another session
- Cursor:pin S – re-execute of same cursor
- Cursor:pin S wait on X – resource is held in X mode by another session
- Library cache: mutex X – Bind variable issues
- Library cache: mutex S – Bind variable issues
- Less costly than latches



# What to Do?

- Share cursors (avoid hard parsing)
  - BIND VARIABLES!!!!!!
  - Cursor\_sharing
- Avoid soft parsing
  - Cursor\_space\_for\_time
  - Session\_cached\_cursors
- Avoid invalidations and reloads
  - Make sure shared pool is large enough



# What Next?

- Determine wait events of concern
- Drill down to specific sections of report for deeper analysis
- Use custom scripts, ADDM and Ash to investigate issues



# Classes

Wait Class

DB/Inst: Snaps: 84084-84108

-> s - second  
-> cs - centisecond - 100th of a second  
-> ms - millisecond - 1000th of a second  
-> us - microsecond - 1000000th of a second  
-> ordered by wait time desc, waits desc

| Wait Class     | Waits       | %Time<br>-outs | Total Wait<br>Time (s) | Avg<br>wait<br>(ms) | Waits<br>/txn |
|----------------|-------------|----------------|------------------------|---------------------|---------------|
| Other          | 153,619,985 | 16.5           | 192,921                | 1                   | 102.3         |
| Concurrency    | 2,536,362   | 26.9           | 128,816                | 51                  | 1.7           |
| User I/O       | 30,594,385  | .0             | 124,207                | 4                   | 20.4          |
| System I/O     | 5,104,873   | .0             | 17,633                 | 3                   | 3.4           |
| Application    | 65,645      | 5.0            | 6,508                  | 99                  | 0.0           |
| Commit         | 267,317     | .0             | 4,234                  | 16                  | 0.2           |
| Configuration  | 553,825     | 69.5           | 858                    | 2                   | 0.4           |
| Network        | 13,513,847  | .0             | 274                    | 0                   | 9.0           |
| Administrative | 30          | 70.0           | 0                      | 10                  | 0.0           |





# Operating System Statistics

Operating System Statistics

DB/Inst: Snaps: 84084-84108

| Statistic              | Total           |
|------------------------|-----------------|
| BUSY_TIME              | 45,601,415      |
| IDLE_TIME              | 6,316,939       |
| IOWAIT_TIME            | 567,343         |
| NICE_TIME              | 810,986         |
| SYS_TIME               | 3,169,946       |
| USER_TIME              | 41,265,848      |
| LOAD                   | 50              |
| RSRC_MGR_CPU_WAIT_TIME | 0               |
| PHYSICAL_MEMORY_BYTES  | 270,208,987,136 |
| NUM_CPUS               | 24              |
| NUM_CPU_SOCKETS        | 4               |



# SQL Areas

**SQL ordered by CPU Time - Sorting, bad paths**

**SQL ordered by Gets - Excessive logical IO**

**SQL ordered by Reads - Cache starvation**

**SQL ordered by Parse Calls - Cursor sharing, cursor caching**

**SQL ordered by Version Count - Versioning is usually due to a bug, check with support**

- **Tune SQL that appears in more than one of these areas**

- **Tune SQL at the top of these sections**



# System Statistics

- Many-many statistics
- Many are not useful to the DBA
- Many are useless for memory area tuning
- Many give ideas of how memory is used, but not how to tune it
- Usually these will confirm what you have already found



# System Statistics

Instance Activity Stats

DB/Inst: Snaps: 84084-84108

| Statistic                        | Total         | per Second | per Trans |
|----------------------------------|---------------|------------|-----------|
| dirty buffers inspected          | 686,267       | 31.8       | 0.5       |
| execute count                    | 78,907,090    | 3,656.2    | 52.6      |
| free buffer inspected            | 161,591,258   | 7,487.4    | 107.6     |
| free buffer requested            | 176,367,274   | 8,172.1    | 117.5     |
| hot buffers moved to head of LRU | 15,346,759    | 711.1      | 10.2      |
| immediate (CR) block cleanout ap | 2,267,512     | 105.1      | 1.5       |
| immediate (CURRENT) block cleano | 5,139,016     | 238.1      | 3.4       |
| no buffer to keep pinned count   | 136,849       | 6.3        | 0.1       |
| no work - consistent read gets   | 4,459,140,613 | 206,616.1  | 2,969.8   |
| opened cursors cumulative        | 26,795,933    | 1,241.6    | 17.9      |
| parse count (failures)           | 160           | 0.0        | 0.0       |
| parse count (hard)               | 398,147       | 18.5       | 0.3       |
| parse count (total)              | 20,200,501    | 936.0      | 13.5      |
| parse time cpu                   | 3,883,178     | 179.9      | 2.6       |
| parse time elapsed               | 7,474,786     | 346.4      | 5.0       |
| physical read total IO requests  | 37,303,810    | 1,728.5    | 24.8      |
| physical reads direct temporary  | 58,378,313    | 2,705.0    | 38.9      |
| physical write total IO requests | 13,738,098    | 636.6      | 9.2       |
| physical writes direct temporary | 58,440,795    | 2,707.9    | 38.9      |



# System Statistics

|                                  |                 |             |          |
|----------------------------------|-----------------|-------------|----------|
| pinned buffers inspected         | 120,843         | 5.6         | 0.1      |
| recursive calls                  | 749,184,714     | 34,713.8    | 499.0    |
| recursive cpu usage              | 39,323,240      | 1,822.1     | 26.2     |
| redo log space requests          | 190             | 0.0         | 0.0      |
| redo log space wait time         | 333             | 0.0         | 0.0      |
| redo synch time                  | 433,625         | 20.1        | 0.3      |
| redo synch writes                | 236,148         | 10.9        | 0.2      |
| redo write time                  | 567,670         | 26.3        | 0.4      |
| redo writer latching time        | 56,827          | 2.6         | 0.0      |
| redo writes                      | 1,127,300       | 52.2        | 0.8      |
| rollback changes - undo records  | 1,395,329       | 64.7        | 0.9      |
| rollbacks only - consistent read | 346,504         | 16.1        | 0.2      |
| session cursor cache hits        | 21,520,355      | 997.2       | 14.3     |
| session logical reads            | 10,474,545,504  | 485,342.3   | 6,976.0  |
| sorts (disk)                     | 3,529           | 0.2         | 0.0      |
| sorts (memory)                   | 9,012,270       | 417.6       | 6.0      |
| sorts (rows)                     | 110,063,794,220 | 5,099,850.2 | 73,302.3 |
| sql area evicted                 | 327,084         | 15.2        | 0.2      |
| sql area purged                  | 29,720          | 1.4         | 0.0      |
| table scans (long tables)        | 1,149,945       | 53.3        | 0.8      |
| table scans (short tables)       | 7,528,140       | 348.8       | 5.0      |
| transaction rollbacks            | 252,407         | 11.7        | 0.2      |



# System Statistics

|                                 |            |       |     |
|---------------------------------|------------|-------|-----|
| user I/O wait time              | 12,422,069 | 575.6 | 8.3 |
| user calls                      | 8,038,839  | 372.5 | 5.4 |
| user commits                    | 1,439,821  | 66.7  | 1.0 |
| user rollbacks                  | 61,684     | 2.9   | 0.0 |
| workarea executions - multipass | 0          | 0.0   | 0.0 |
| workarea executions - onepass   | 5,293      | 0.3   | 0.0 |
| workarea executions - optimal   | 7,113,060  | 329.6 | 4.7 |



# Instance Activity Statistics

Instance Activity Stats - Absolute

-> Statistics with absolute values (should not be diffed)

| Statistic                  | Begin Value | End Value  |
|----------------------------|-------------|------------|
| -----                      | -----       | -----      |
| session cursor cache count | 28,024,069  | 28,789,659 |
| opened cursors current     | 2,921       | 6,982      |
| workarea memory allocated  | 289,532     | 2,531,741  |
| logons current             | 144         | 287        |
| -----                      | -----       | -----      |



# Tablespace/File IO Reports

- Helps confirm IO issues
- Also helps with temp area issue determination





# Tablespace IO

Tablespace IO Stats

-> ordered by IOs (Reads + Writes) desc

Tablespace

| Tablespace              | Reads      | AV Reads/s | AV Rd(ms) | AV Blks/Rd | Writes    | AV Writes/s | Buffer Waits | AV Buf Wt(ms) |
|-------------------------|------------|------------|-----------|------------|-----------|-------------|--------------|---------------|
| TEMP                    | 11,484,000 | 532        | 16.3      | 4.1        | 3,478,365 | 161         | 12,266       | 2.0           |
| REPMAN_TEMP             | 1,703,767  | 79         | 27.2      | 8.2        | 1,457,241 | 68          | 0            | 0.0           |
| UNDOTBS3                | 30,012     | 1          | 8.0       | 1.0        | 1,512,571 | 70          | 142,889      | 1.1           |
| RSNET_DTSA              | 1,496,441  | 69         | 1.2       | 2.0        | 2,454     | 0           | 130,753      | 1.3           |
| LOREAL_D_CVS_DAILY_ITSA | 846,665    | 39         | 0.9       | 1.0        | 338       | 0           | 0            | 0.0           |



# Buffer Pool Statistics

Buffer Pool Statistics DB/Inst: CC1/cc1 Snaps: 84084-84108

-> Standard block size Pools D: default, K: keep, R: recycle

-> Default Pools for other block sizes: 2k, 4k, 8k, 16k, 32k

| P | Number of Pool<br>Buffers | Hit% | Buffer<br>Gets | Physical<br>Reads | Physical<br>Writes | Free Buff<br>Wait | Writ<br>Comp<br>Wait | Buffer<br>Busy<br>Waits |
|---|---------------------------|------|----------------|-------------------|--------------------|-------------------|----------------------|-------------------------|
| D | 3,361,107                 | 96   | 3,643,978,600  | 163,055,679       | 14,338,623         | 0                 | 0                    | 711,281                 |
| K | 321,600                   | 100  | 2,527,600,634  | 7,379             | 28,755             | 0                 | 0                    | 123                     |

- Note that there are Buffer Busy Waits, but no Free Buffer Waits
- These are due to hot block contention
- Increasing memory probably won't help with this
- However...also look at db file sequential read waits and the cache advisory section



# Buffer Pool Advisory Section

## Buffer Pool Advisory

-> Only rows with estimated physical reads >0 are displayed

-> ordered by Block Size, Buffers For Estimate

| P   | Size for<br>Est (M) | Size<br>Factor | Buffers for<br>Estimate | Est<br>Phys<br>Read<br>Factor | Estimated<br>Physical Reads |
|-----|---------------------|----------------|-------------------------|-------------------------------|-----------------------------|
| --- | -----               | -----          | -----                   | -----                         | -----                       |
| D   | 5,344               | .1             | 335,670                 | 1.9                           | 15,767,325,073              |
| D   | 10,688              | .2             | 671,340                 | 1.4                           | 11,371,357,960              |
| ... |                     |                |                         |                               |                             |
| D   | 106,880             | 2.0            | 6,713,400               | 1.0                           | 7,964,367,701               |
| K   | 512                 | .1             | 32,160                  | 102.8                         | 3,507,100,178               |
| K   | 1,024               | .2             | 64,320                  | 7.8                           | 264,615,629                 |
| K   | 1,536               | .3             | 96,480                  | 1.4                           | 49,384,590                  |
| ... |                     |                |                         |                               |                             |
| K   | 10,240              | 2.0            | 643,200                 | 1.0                           | 32,639,643                  |

-----



# Buffer Pool Advisory Section

- As you can see, this report shows even doubling the default or keep would be no benefit
- Let's look at one that would benefit from increased buffer pool



# Buffer Pool Advisor Section

|                                         |                 |                   |             |                |                    |                |                |                   |
|-----------------------------------------|-----------------|-------------------|-------------|----------------|--------------------|----------------|----------------|-------------------|
| Buffer Pool Statistics                  |                 |                   | DB/Inst:    |                | Snaps: 26064-26097 |                |                |                   |
| -> Standard block size Pools            |                 |                   | D: default, |                | K: keep,           |                | R: recycle     |                   |
| -> Default Pools for other block sizes: |                 |                   | 2k,         |                | 4k,                |                | 8k, 16k, 32k   |                   |
| P                                       | Number of Pools | Pool Buffers Hit% | Buffer Gets | Physical Reads | Physical Writes    | Free Buff Wait | Writ Comp Wait | Buffer Busy Waits |
| D                                       | 818,201         | 99                | 130,795,544 | 1,578,580      | 276,075            | 0              | 0              | 3,418             |

- Before we go there...look here
- Notice no free buffer waits



# Buffer Pool Advisor Section

Buffer Pool Advisory

Snap: 26097

-> Only rows with estimated physical reads >0 are displayed

-> ordered by Block Size, Buffers For Estimate

| P   | Size for<br>Est (M) | Size<br>Factor | Buffers for<br>Estimate | Est<br>Phys<br>Read<br>Factor | Estimated<br>Physical Reads |
|-----|---------------------|----------------|-------------------------|-------------------------------|-----------------------------|
| --- | -----               | -----          | -----                   | -----                         | -----                       |
| D   | 656                 | .1             | 81,139                  | 2.0                           | 125,592,784                 |
| D   | 1,312               | .2             | 162,278                 | 1.8                           | 113,080,052                 |
| ... |                     |                |                         |                               |                             |
| D   | 10,496              | 1.6            | 1,298,224               | 0.7                           | 42,942,366                  |
| D   | 11,152              | 1.7            | 1,379,363               | 0.7                           | 41,649,501                  |
| D   | 11,808              | 1.8            | 1,460,502               | 0.6                           | 40,403,058                  |
| --- | -----               | -----          | -----                   | -----                         | -----                       |

• Notice that at 1.8 times the current size physical reads down by 40%

• Db file sequential reads was 13% of waits



# PGA Analysis

- Several of the next sections deal with PGA
- `PGA_AGGREGATE_TARGET` sets the PGA area
- 5% of `PGA_AGGREGATE_TARGET` can be allocated to each session up to a maximum of “`_PGA_MAX_SIZE`” which is usually 200 or 500 megabytes
- Manually setting `SORT_AREA_SIZE` or `HASH_AREA_SIZE` overrides at the session level
- Some processes such as RMAN and shared servers don't use `PGA_AGGREGATE_TARGET` but use the old manual settings, indicated by 4-8 or 8-16m sorts even with adequate `PGA_AGGREGATE_TARGET` settings



# PGA Analysis

PGA Aggr Summary

DB/Inst: Snaps: 84084-84108

-> PGA cache hit % - percentage of W/A (WorkArea) data processed only in-memory

| PGA Cache Hit % | W/A MB Processed | Extra W/A MB Read/Written |
|-----------------|------------------|---------------------------|
|-----------------|------------------|---------------------------|

|      |           |         |
|------|-----------|---------|
| 82.1 | 4,495,435 | 979,073 |
|------|-----------|---------|





# PGA Analysis

PGA Aggr Target Stats

DB/Inst: Snaps: 84084-84108

- > B: Begin snap E: End snap (rows identified with B or E contain data which is absolute i.e. not diffed over the interval)
- > Auto PGA Target - actual workarea memory target
- > W/A PGA Used - amount of memory used for all Workareas (manual + auto)
- > %PGA W/A Mem - percentage of PGA memory allocated to workareas
- > %Auto W/A Mem - percentage of workarea memory controlled by Auto Mem Mgmt
- > %Man W/A Mem - percentage of workarea memory under manual control

|   | PGA Aggr<br>Target(M) | Auto PGA<br>Target(M) | PGA Mem<br>Alloc(M) | W/A PGA<br>Used(M) | %PGA<br>W/A<br>Mem | %Auto<br>W/A<br>Mem | %Man<br>W/A<br>Mem | Global Mem<br>Bound(K) |
|---|-----------------------|-----------------------|---------------------|--------------------|--------------------|---------------------|--------------------|------------------------|
| B | 5,120                 | 4,320                 | 1,680.5             | 193.5              | 11.5               | 99.7                | .3                 | 524,280                |
| E | 5,120                 | 4,202                 | 4,400.5             | 2,219.2            | 50.4               | 99.9                | .1                 | 524,280                |



# PGA Analysis

PGA Aggr Target Histogram

Snapshots: 84084-84108

| Low     | High    | Total Execs | Optimal Execs | 1-Pass Execs | M-Pass Execs |
|---------|---------|-------------|---------------|--------------|--------------|
| Optimal | Optimal |             |               |              |              |
| 2K      | 4K      | 6,308,435   | 6,308,435     | 0            | 0            |
| 64K     | 128K    | 32,500      | 32,500        | 0            | 0            |
| 128K    | 256K    | 36,535      | 36,535        | 0            | 0            |
| 256K    | 512K    | 47,477      | 47,477        | 0            | 0            |
| 512K    | 1024K   | 353,344     | 353,344       | 0            | 0            |
| 1M      | 2M      | 201,558     | 201,558       | 0            | 0            |
| 2M      | 4M      | 22,468      | 22,468        | 0            | 0            |
| 4M      | 8M      | 21,796      | 21,725        | 71           | 0            |
| 8M      | 16M     | 28,892      | 28,714        | 178          | 0            |
| 16M     | 32M     | 30,478      | 30,346        | 132          | 0            |
| 32M     | 64M     | 19,898      | 18,690        | 1,208        | 0            |
| 64M     | 128M    | 9,080       | 7,284         | 1,796        | 0            |
| 128M    | 256M    | 1,682       | 732           | 950          | 0            |
| 256M    | 512M    | 734         | 179           | 555          | 0            |
| 512M    | 1024M   | 329         | 58            | 271          | 0            |
| 1G      | 2G      | 131         | 14            | 117          | 0            |
| 2G      | 4G      | 17          | 4             | 13           | 0            |



# PGA Analysis

- $SS * 100 / 2.5 * NOS$ 
  - SS- sort size
  - NOS - number of sorts
- $128M * 100 / 2.5 = 5120$
- Hash gets full 5%
- Sort gets  $\frac{1}{2}$  of hash
- $\frac{1}{2}$  of 5 is 2.5
- 10|10 rule:
  - 10% of sessions are active
  - 10% of active sessions doing sorts
- 10% of 208 sessions is 20.8; 10% of 20.8 is 2
- $2 * 5120$  is 10420 gb
- Let's see what the advisor says...



# PGA Analysis

PGA Memory Advisory

Snap: 84108

-> When using Auto Memory Mgmt, minimally choose a pga\_aggregate\_target value where Estd PGA Overalloc Count is 0

| PGA Target<br>Est (MB) | Size<br>Factr | W/A MB<br>Processed | Estd Extra<br>W/A MB Read/<br>Written to Disk | Estd PGA<br>Cache<br>Hit % | Estd PGA<br>Overalloc<br>Count |
|------------------------|---------------|---------------------|-----------------------------------------------|----------------------------|--------------------------------|
| 640                    | 0.1           | 167,211,258.2       | 133,135,735.1                                 | 56.0                       | 18,399                         |
| 1,280                  | 0.3           | 167,211,258.2       | 65,323,839.1                                  | 72.0                       | 1,933                          |
| 2,560                  | 0.5           | 167,211,258.2       | 34,248,434.3                                  | 83.0                       | 1,842                          |
| 3,840                  | 0.8           | 167,211,258.2       | 28,768,442.5                                  | 85.0                       | 1,803                          |
| 5,120                  | 1.0           | 167,211,258.2       | 19,597,963.1                                  | 90.0                       | 860                            |
| 6,144                  | 1.2           | 167,211,258.2       | 14,425,059.9                                  | 92.0                       | 845                            |
| 7,168                  | 1.4           | 167,211,258.2       | 13,624,988.5                                  | 92.0                       | 839                            |
| 8,192                  | 1.6           | 167,211,258.2       | 12,882,080.7                                  | 93.0                       | 0                              |
| 9,216                  | 1.8           | 167,211,258.2       | 12,146,846.2                                  | 93.0                       | 0                              |
| 10,240                 | 2.0           | 167,211,258.2       | 11,689,119.6                                  | 93.0                       | 0                              |
| 15,360                 | 3.0           | 167,211,258.2       | 10,710,132.3                                  | 94.0                       | 0                              |
| 20,480                 | 4.0           | 167,211,258.2       | 10,563,396.2                                  | 94.0                       | 0                              |
| 30,720                 | 6.0           | 167,211,258.2       | 10,539,291.8                                  | 94.0                       | 0                              |
| 40,960                 | 8.0           | 167,211,258.2       | 10,539,291.8                                  | 94.0                       | 0                              |



# PGA Analysis

- Advisor says 8 gb
- Calculation says 10 gb
- Take whichever you feel will give best results
- Won't be used unless it needs to be



# Shared Pool Advisor

- Use reloads and used percentages to guide you
- The advisor rarely has any meaningful information
- Even when reloads are huge and other factors show the pool should be increased, or, decreased it has told me things where fine



# Shared Pool Advisor

Shared Pool Advisory

DB/Inst: Snap: 84108

-> SP: Shared Pool      Est LC: Estimated Library Cache      Factr: Factor

-> Note there is often a 1:Many correlation between a single logical object in the Library Cache, and the physical number of memory objects associated with it. Therefore comparing the number of Lib Cache objects (e.g. in v\$librarycache), with the number of Lib Cache Memory Objects is invalid.

| Shared<br>Pool<br>Size(M) | SP<br>Size<br>Factr | Est LC<br>Size<br>(M) | Est LC<br>Mem Obj | Est LC<br>Time<br>Saved<br>(s) | Est LC<br>Time<br>Saved<br>Factr | Est LC<br>Load<br>Time<br>(s) | Est LC<br>Load<br>Time<br>Factr | Est LC<br>Mem<br>Obj Hits |
|---------------------------|---------------------|-----------------------|-------------------|--------------------------------|----------------------------------|-------------------------------|---------------------------------|---------------------------|
| 2,160                     | .4                  | 619                   | 76,837            | #####                          | 1.0                              | #####                         | 2.4                             | 88,538,740                |
| 2,736                     | .5                  | 1,188                 | 95,749            | #####                          | 1.0                              | #####                         | 2.1                             | 88,936,333                |
| 3,312                     | .6                  | 1,761                 | 110,785           | #####                          | 1.0                              | #####                         | 1.8                             | 89,297,339                |
| 3,888                     | .7                  | 2,333                 | 125,755           | #####                          | 1.0                              | #####                         | 1.6                             | 89,610,155                |
| ...                       |                     |                       |                   |                                |                                  |                               |                                 |                           |
| 8,496                     | 1.5                 | 6,916                 | 238,592           | #####                          | 1.0                              | #####                         | .5                              | 91,076,008                |
| 9,072                     | 1.6                 | 7,489                 | 252,248           | #####                          | 1.0                              | #####                         | .4                              | 91,187,541                |
| 9,648                     | 1.7                 | 8,061                 | 264,748           | #####                          | 1.0                              | #####                         | .3                              | 91,291,114                |
| 10,224                    | 1.8                 | 8,632                 | 274,837           | #####                          | 1.0                              | #####                         | .3                              | 91,388,340                |
| 10,800                    | 1.9                 | 9,201                 | 284,723           | #####                          | 1.0                              | #####                         | .2                              | 91,480,577                |
| 11,376                    | 2.0                 | 9,774                 | 293,073           | #####                          | 1.0                              | #####                         | .2                              | 91,569,191                |



# SGA Target Advisor

- I believe you should set the base parameters and then allow SGA\_TARGET to set itself
- Other calculation schemes involve adding up the proposed sizes needed and then setting the value
- Same logic applies to MEMORY\_TARGET only include PGA\_AGGREGATE\_TARGET in the base
- Then set SGA\_MAX\_SIZE or MEMORY\_MAX\_SIZE 10-20% higher than their associated TARGET values.





# SGA Target

SGA Target Advisory

Snap: 26097

| SGA Target<br>Size (M) | SGA Size<br>Factor | Est DB<br>Time (s) | Est Physical<br>Reads |
|------------------------|--------------------|--------------------|-----------------------|
| 2,048                  | 0.3                | 8,957,297          | 125,595,242           |
| 4,096                  | 0.5                | 745,014            | 99,011,933            |
| 6,144                  | 0.8                | 683,241            | 84,201,232            |
| 8,192                  | 1.0                | 603,250            | 62,255,994            |
| 10,240                 | 1.3                | 573,691            | 48,902,083            |
| 12,288                 | 1.5                | 573,630            | 42,944,185            |
| 14,336                 | 1.8                | 573,630            | 37,745,809            |
| 16,384                 | 2.0                | 573,630            | 37,745,809            |

- At 2x shows a 60% reduction in PR, same as doubling the cache for this instance



# Streams Pool Advisor

- Only valid if you use streams
- If it shows you have spills to disk make pool larger



# Streams Pool Advisor

Streams Pool Advisory

DB/Inst: Snap: 84108

| Size for<br>Est (MB) | Size<br>Factor | Est Spill<br>Count | Est Spill<br>Time (s) | Est Unspill<br>Count | Est Unspill<br>Time (s) |
|----------------------|----------------|--------------------|-----------------------|----------------------|-------------------------|
| 16                   | 0.5            | 0                  | 0                     | 0                    | 0                       |
| 32                   | 1.0            | 0                  | 0                     | 0                    | 0                       |
| 48                   | 1.5            | 0                  | 0                     | 0                    | 0                       |
| 64                   | 2.0            | 0                  | 0                     | 0                    | 0                       |
| 80                   | 2.5            | 0                  | 0                     | 0                    | 0                       |
| ...                  |                |                    |                       |                      |                         |
| 240                  | 7.5            | 0                  | 0                     | 0                    | 0                       |
| 256                  | 8.0            | 0                  | 0                     | 0                    | 0                       |
| 272                  | 8.5            | 0                  | 0                     | 0                    | 0                       |
| 288                  | 9.0            | 0                  | 0                     | 0                    | 0                       |
| 304                  | 9.5            | 0                  | 0                     | 0                    | 0                       |
| 320                  | 10.0           | 0                  | 0                     | 0                    | 0                       |



# Java Pool Advisor

- Another area usually not used
- If you use it you will get errors if it is too small
- Under AMM will grow but usually won't shrink



# Java Pool Advisor

Java Pool Advisory

DB/Inst: Snap: 37

| Java<br>Pool<br>Size(M) | JP<br>Size<br>Factr | Est LC<br>Size<br>(M) | Est LC<br>Mem Obj | Est LC<br>Time<br>Saved<br>(s) | Est LC<br>Time<br>Saved<br>Factr | Est LC<br>Load<br>Time<br>(s) | Est LC<br>Load<br>Time<br>Factr | Est LC<br>Mem<br>Obj Hits |
|-------------------------|---------------------|-----------------------|-------------------|--------------------------------|----------------------------------|-------------------------------|---------------------------------|---------------------------|
| 64                      | .5                  | 10                    | 168               | 10                             | 1.0                              | 11,974                        | 1.0                             | 389                       |
| 128                     | 1.0                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |
| 192                     | 1.5                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |
| 256                     | 2.0                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |



# Buffer Wait Analysis

- Look here to see what is causing buffer waits
- Data blocks usually predominate
- High buffer waits for data blocks plus high db file sequential reads will usually indicate memory starvation even without free buffer waits



# Buffer Wait Analysis

Buffer Wait Statistics

DB/Inst:

Snap: 84084-84108

-> ordered by wait time desc, waits desc

| Class             | Waits   | Total Wait Time (s) | Avg Time (ms) |
|-------------------|---------|---------------------|---------------|
| data block        | 539,313 | 1,770               | 3             |
| 1st level bmb     | 5,873   | 88                  | 15            |
| undo block        | 49,801  | 71                  | 1             |
| undo header       | 35,848  | 66                  | 2             |
| file header block | 57,833  | 39                  | 1             |
| segment header    | 12,980  | 28                  | 2             |
| 3rd level bmb     | 4,900   | 22                  | 5             |
| 2nd level bmb     | 5,206   | 20                  | 4             |
| extent map        | 73      | 0                   | 5             |



# Enqueues and Latches

- Enqueues are usually for physical objects or transactions and rarely memory related
- Latches can be used to verify findings





# Latches

## Latch Activity

DB/Inst: Snaps: 84084-84108

- > "Get Requests", "Pct Get Miss" and "Avg Slps/Miss" are statistics for willing-to-wait latch get requests
- > "NoWait Requests", "Pct NoWait Miss" are for no-wait latch get requests
- > "Pct Misses" for both should be very close to 0.0

| Latch Name               | Get Requests  | Pct Get Miss | Avg Slps /Miss | Wait Time (s) | NoWait Requests | Pct NoWait Miss |
|--------------------------|---------------|--------------|----------------|---------------|-----------------|-----------------|
| SGA IO buffer pool latch | 154,979       | 0.0          | 0.7            | 0             | 261,675         | 0.3             |
| SQL memory manager latch | 21,353        | 12.7         | 1.2            | 32            | 7,066           | 0.1             |
| active service list      | 2,524,503     | 1.7          | 0.2            | 10            | 7,615           | 0.1             |
| cache buffers lru chain  | 66,153,522    | 1.0          | 0.6            | 1172          | 469,078,119     | 1.8             |
| cache table scan latch   | 3,749,595     | 0.3          | 0.3            | 9             | 3,986,131       | 0.3             |
| library cache            | 272,976,941   | 0.1          | 0.4            | 2242          | 1,582,568       | 216.3           |
| library cache lock       | 179,441,868   | 0.1          | 0.1            | 34            | 58,383          | 0.1             |
| object queue header oper | 883,815,615   | 0.1          | 0.1            | 365           | 45,167          | 0.2             |
| process queue reference  | 8,224,189,548 | 0.0          | 0.0            | 15            | 613,910,934     | 4.3             |
| redo allocation          | 8,482,485     | 8.1          | 0.1            | 271           | 435,767,620     | 1.4             |
| row cache objects        | 871,266,120   | 0.4          | 0.2            | 1765          | 580,035         | 1.6             |
| simulator lru latch      | 7,890         | 1.1          | 1.1            | 1             | #####           | 17.1            |
| undo global data         | 68,788,574    | 0.0          | 0.1            | 8             | 42,419          | 0.1             |



# Latches

- When a latch can't be obtained a process will "sleep"
- The number of CPU cycles is set by the "\_SPIN\_COUNT" setting
- Defaults to 2000, was set in the era of slow CPUs
- Experts agree setting as high as 8000-10000 is acceptable
- Reduces sleeps and improved performance



# Latch Sleeps

Latch Sleep Breakdown

DB/Inst: Snaps: 84084-84108

-> ordered by misses desc

Latch Name

| Get Requests                  | Misses     | Sleeps  | Spin Gets | Sleep1 | Sleep2 | Sleep3 |
|-------------------------------|------------|---------|-----------|--------|--------|--------|
| cache buffers chains          |            |         |           |        |        |        |
| 9,561,948,628                 | #####      | 402,582 | #####     | 0      | 0      | 0      |
| session allocation            |            |         |           |        |        |        |
| 570,516,196                   | 14,728,230 | 461,868 | #####     | 0      | 0      | 0      |
| messages                      |            |         |           |        |        |        |
| 105,586,860                   | 3,521,238  | 16,783  | 3,506,276 | 0      | 0      | 0      |
| row cache objects             |            |         |           |        |        |        |
| 871,266,120                   | 3,508,138  | 655,697 | 2,914,340 | 0      | 0      | 0      |
| process queue reference       |            |         |           |        |        |        |
| 8,224,189,548                 | 2,645,262  | 19,802  | 2,629,967 | 0      | 0      | 0      |
| active checkpoint queue latch |            |         |           |        |        |        |
| 50,214,009                    | 1,082,455  | 13,133  | 1,070,658 | 0      | 0      | 0      |
| redo writing                  |            |         |           |        |        |        |
| 52,201,162                    | 1,026,204  | 6,893   | 1,020,132 | 0      | 0      | 0      |
| object queue header operation |            |         |           |        |        |        |
| 883,815,615                   | 936,882    | 137,869 | 812,882   | 0      | 0      | 0      |



# Latch Miss Sources

Latch Miss Sources

DB/Inst: Snaps: 84084-84108

-> only latches with sleeps are shown

-> ordered by name, sleeps desc

| Latch Name               | Where                      | NoWait<br>Misses | Sleeps  | Waiter<br>Sleeps |
|--------------------------|----------------------------|------------------|---------|------------------|
| In memory undo latch     | ktiFlush: child            | 0                | 12,727  | 3,775            |
| ...                      |                            |                  |         |                  |
| cache buffers chains     | kcbgtcr: kslbegin excl     | 0                | 821,219 | 689,229          |
| ...                      |                            |                  |         |                  |
| cache buffers lru chain  | kcbbxsv: move to being wri | 0                | 19,728  | 87,892           |
| ...                      |                            |                  |         |                  |
| library cache            | kglpndl: child: after proc | 0                | 28,574  | 10,336           |
| ...                      |                            |                  |         |                  |
| library cache lock       | kgllkdl: child: cleanup    | 0                | 7,106   | 6,749            |
| ...                      |                            |                  |         |                  |
| object queue header oper | kcbw_unlink_q              | 0                | 139,971 | 27,436           |
| ...                      |                            |                  |         |                  |
| parameter table allocati | ksp_param_table_free       | 0                | 60,653  | 60,605           |
| row cache objects        | kqrpre: find obj           | 0                | 381,022 | 379,948          |
| ...                      |                            |                  |         |                  |
| session allocation       | ksuprc                     | 0                | 191,233 | 175,384          |
| ...                      |                            |                  |         |                  |
| shared pool              | kghalo                     | 0                | 194,517 | 123,739          |
| ...                      |                            |                  |         |                  |
| simulator hash latch     | <u>kcbsacc: lookup dba</u> | 0                | 16,413  | 26,348           |



# Dictionary Cache

- Misses can be costly
- Only correction is larger shared pool
- Used to be individually controlled by DC parameters (v6)
- Now controlled automatically



# Dictionary Cache

Dictionary Cache Stats

DB/Inst: Snaps: 84084-84108

-> "Pct Misses" should be very low (< 2% in most cases)

-> "Final Usage" is the number of cache entries being used

| Cache                | Get<br>Requests | Pct<br>Miss | Scan<br>Reqs | Pct<br>Miss | Mod<br>Reqs | Final<br>Usage |
|----------------------|-----------------|-------------|--------------|-------------|-------------|----------------|
| dc_awr_control       | 442             | 9.3         | 0            | N/A         | 48          | 1              |
| dc_constraints       | 60,034          | 41.5        | 0            | N/A         | 60,034      | 49             |
| dc_files             | 559             | 97.3        | 0            | N/A         | 0           | 0              |
| dc_global_oids       | 20,580          | 3.0         | 0            | N/A         | 0           | 25             |
| dc_histogram_data    | 7,832,289       | 2.0         | 0            | N/A         | 19,621      | 7,472          |
| dc_histogram_defs    | 5,795,619       | 15.5        | 0            | N/A         | 322,942     | 6,953          |
| dc_object_grants     | 99,052          | 11.6        | 0            | N/A         | 0           | 370            |
| dc_object_ids        | 105,285,794     | 2.3         | 0            | N/A         | 38,220      | 66,669         |
| dc_objects           | 53,343,945      | 4.8         | 0            | N/A         | 161,568     | 8,250          |
| dc_profiles          | 60,308          | 0.1         | 0            | N/A         | 0           | 5              |
| dc_segments          | 19,064,933      | 58.4        | 0            | N/A         | 256,828     | 16,244         |
| dc_sequences         | 137,407         | 0.9         | 0            | N/A         | 137,407     | 35             |
| dc_table_scns        | 180,866         | 4.2         | 0            | N/A         | 88          | 14             |
| dc_tablespace_quotas | 101,823         | 2.5         | 0            | N/A         | 101,679     | 43             |
| dc tablespaces       | 11,521,250      | 1.3         | 0            | N/A         | 18          | 10,945         |
| dc_usernames         | 750,849         | 0.4         | 0            | N/A         | 0           | 104            |
| global database name | 118             | 6.8         | 0            | N/A         | 0           | 1              |
| kqlsubheap_object    | 1               | 100.0       | 0            | N/A         | 0           | 0              |
| outstanding_alerts   | 394,019         | 41.7        | 0            | N/A         | 0           | 10,943         |



# Library Cache Activity

- Usually only see issues in SQL and PL/SQL areas
- Invalidations high means DDL activity
- Reloads high means SQL or PL/SQL issues



# Library Cache Activity

Library Cache Activity

DB/Inst: Snaps: 84084-84108

-> "Pct Misses" should be very low

| Namespace       | Get<br>Requests | Pct<br>Miss | Pin<br>Requests | Pct<br>Miss | Reloads   | Invali-<br>dations |
|-----------------|-----------------|-------------|-----------------|-------------|-----------|--------------------|
| BODY            | 106,197         | 0.3         | 2,372,306       | 0.1         | 1,288     | 0                  |
| CLUSTER         | 11,232          | 0.4         | 19,082          | 0.7         | 81        | 0                  |
| INDEX           | 1,728,015       | 1.9         | 28,528,015      | 0.2         | 17,973    | 0                  |
| PIPE            | 9               | 0.0         | 9               | 0.0         | 0         | 0                  |
| SQL AREA        | 3,482,631       | 14.7        | 108,055,227     | 2.4         | 1,276,544 | #####              |
| TABLE/PROCEDURE | 290,691         | 6.1         | 32,986,432      | 0.5         | 76,695    | 0                  |
| TRIGGER         | 621,821         | 0.1         | 2,312,915       | 0.1         | 2,423     | 0                  |





# Process Statistics

## Process Memory Summary

DB/Inst: Snaps: 84084-84108

- > B: Begin snap E: End snap
- > All rows below contain absolute values (i.e. not diffed over the interval)
- > Max Alloc is Maximum PGA Allocation size at snapshot time
- > Hist Max Alloc is the Historical Max Allocation for still-connected processes
- > ordered by Begin/End snapshot, Alloc (MB) desc

|            |         |         |       |         |       | Hist  |      |       |
|------------|---------|---------|-------|---------|-------|-------|------|-------|
|            | Alloc   | Used    | Avg   | Std Dev | Max   | Max   |      |       |
| Category   | (MB)    | (MB)    | Alloc | Alloc   | Alloc | Alloc | Num  | Num   |
|            |         |         | (MB)  | (MB)    | (MB)  | (MB)  | Proc | Alloc |
| B Freeable | 981.6   | .0      | 6.2   | 23.0    | 141   | N/A   | 158  | 158   |
| Other      | 647.4   | N/A     | 3.3   | 5.8     | 30    | 102   | 195  | 191   |
| PL/SQL     | 38.8    | 7.0     | .2    | .6      | 3     | 3     | 193  | 154   |
| SQL        | 15.7    | 10.9    | .1    | .2      | 2     | 1,289 | 174  | 128   |
| E SQL      | 1,953.6 | 1,939.0 | 6.7   | 36.8    | 488   | 1,289 | 291  | 264   |
| Other      | 1,772.1 | N/A     | 5.6   | 16.7    | 177   | 498   | 315  | 311   |
| Freeable   | 640.8   | .0      | 4.1   | 14.2    | 92    | N/A   | 156  | 156   |
| PL/SQL     | 77.0    | 13.8    | .2    | .5      | 3     | 13    | 313  | 313   |



# SGA Memory Summary

SGA Memory Summary

DB/Inst: Snaps: 84084-84108

| SGA regions      | Begin Size (Bytes) | End Size (Bytes)<br>(if different) |
|------------------|--------------------|------------------------------------|
| -----            | -----              | -----                              |
| Database Buffers | 61,035,511,808     | 61,538,828,288                     |
| Fixed Size       | 2,212,832          |                                    |
| Redo Buffers     | 14,561,280         |                                    |
| Variable Size    | 7,667,190,816      | 7,163,874,336                      |
|                  | -----              |                                    |
| sum              | 68,719,476,736     |                                    |
|                  | -----              |                                    |



# SGA Breakdown

SGA breakdown difference

DB/Inst: Snaps: 84084-84108

-> ordered by Pool, Name

-> N/A value for Begin MB or End MB indicates the size of that Pool/Name was insignificant, or zero in that snapshot

| Pool   | Name                      | Begin MB | End MB   | % Diff  |
|--------|---------------------------|----------|----------|---------|
| java   | free memory               | 160.0    | 160.0    | 0.00    |
| large  | free memory               | 1,022.6  | 1,020.6  | -0.19   |
| shared | CCursor                   | 483.3    | 425.2    | -12.04  |
| shared | Checkpoint queue          | 175.8    | 175.8    | 0.00    |
| shared | KGH: NO ACCESS            | 2,445.0  | 3,079.3  | 25.94   |
| shared | KQR L PO                  | 355.6    | N/A      | -100.00 |
| shared | KQR M PO                  | 430.4    | N/A      | -100.00 |
| shared | PCursor                   | 221.9    | 190.5    | -14.16  |
| shared | db_block_hash_buckets     | 180.0    | 180.0    | 0.00    |
| shared | free memory               | 2,191.3  | 2,890.1  | 31.89   |
| shared | kzctxgjsi ksuseclid memor | 1,644.5  | 1,676.7  | 1.96    |
| shared | library cache             | 334.2    | 308.2    | -7.79   |
| shared | obj stat memo             | 466.8    | 474.8    | 1.71    |
| shared | partitioning d            | 131.8    | N/A      | -100.00 |
| shared | sql area                  | 483.6    | 236.8    | -51.04  |
| stream | free memory               | 31.9     | 31.9     | 0.00    |
|        | buffer_cache              | 58,208.0 | 58,688.0 | 0.82    |
|        | fixed_sga                 | 2.1      | 2.1      | 0.00    |
|        | log_buffer                | 13.9     | 13.9     | 0.00    |



# Resize Operations

- Using AMM resize operations are automatic
- If there is room to grow, shouldn't have to borrow (steal) from other areas
- If `SGA_TARGET=SGA_MAX_SIZE` or `MEMORY_TARGET=MEMORY_MAX_SIZE` then no room to grow



# Resize Operations

Memory Dynamic Components

DB/Inst: Snaps: 22930-22958

-> Min/Max sizes since instance startup

-> Oper Types/Modes: INItializing, GROw, SHRink, STAtic/IMMediate, DEFerred

-> ordered by Component

| Component       | Begin Snap<br>Size (Mb) | Current<br>Size (Mb) | Min<br>Size (Mb) | Max<br>Size (Mb) | Oper<br>Count | Last Op<br>Typ/Mod |
|-----------------|-------------------------|----------------------|------------------|------------------|---------------|--------------------|
| ASM Buffer Cach | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 16K buf | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 2K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 32K buf | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 4K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 8K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT buffer  | 20,480.00               | 20,480.00            | 19,712.00        | 20,736.00        | 0             | GRO/DEF            |
| KEEP buffer cac | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| PGA Target      | 14,848.00               | 14,848.00            | 14,848.00        | 14,848.00        | 0             | STA/               |
| RECYCLE buffer  | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| SGA Target      | 28,160.00               | 28,160.00            | 28,160.00        | 28,160.00        | 0             | STA/               |
| Shared IO Pool  | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| java pool       | 1,024.00                | 1,024.00             | 1,024.00         | 1,024.00         | 0             | SHR/DEF            |
| large pool      | 256.00                  | 256.00               | 256.00           | 256.00           | 0             | STA/               |
| shared pool     | 4,096.00                | 4,096.00             | 3,840.00         | 4,864.00         | 0             | SHR/DEF            |
| streams pool    | 2,048.00                | 2,048.00             | 2,048.00         | 2,048.00         | 0             | STA/               |

**tms** V\$SGA\_RESIZE\_OPS -10g  
V\$MEMORY\_RESIZE\_OPS - 11g

- A dynamic performance view
- Use to get details of resize operations
- Every memory change since startup



# V\$SGA\_RESIZE\_OPS

Date: 02/18/08

Page: 1

Time: 01:04 PM

Component Resize Operations

SYSTEM

ault10g1 database

| COMPONENT            | Oper   | OPER_MODE | INITIAL_SIZE | TARGET_SIZE | FINAL_SIZE | STATUS   | START_TIME | END_TIME   |
|----------------------|--------|-----------|--------------|-------------|------------|----------|------------|------------|
| shared pool          | SHRINK | DEFERRED  | 318767104    | 301989888   | 301989888  | COMPLETE | 0217 15:23 | 0217 15:23 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1040187392   | 1056964608  | 1056964608 | COMPLETE | 0217 15:23 | 0217 15:23 |
| shared pool          | SHRINK | DEFERRED  | 301989888    | 285212672   | 285212672  | COMPLETE | 0217 15:24 | 0217 15:24 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1073741824   | 1090519040  | 1090519040 | COMPLETE | 0217 15:24 | 0217 15:24 |
| shared pool          | SHRINK | DEFERRED  | 285212672    | 268435456   | 268435456  | COMPLETE | 0217 15:24 | 0217 15:24 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1056964608   | 1073741824  | 1073741824 | COMPLETE | 0217 15:24 | 0217 15:24 |
| shared pool          | SHRINK | DEFERRED  | 268435456    | 251658240   | 251658240  | COMPLETE | 0217 15:25 | 0217 15:25 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1090519040   | 1107296256  | 1107296256 | COMPLETE | 0217 15:25 | 0217 15:25 |
| shared pool          | SHRINK | DEFERRED  | 251658240    | 234881024   | 234881024  | COMPLETE | 0217 15:28 | 0217 15:28 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1107296256   | 1124073472  | 1124073472 | COMPLETE | 0217 15:28 | 0217 15:28 |
| shared pool          | SHRINK | DEFERRED  | 234881024    | 218103808   | 218103808  | COMPLETE | 0217 15:28 | 0217 15:28 |



# Parameters

init.ora Parameters

DB/Inst:      Snaps: 84084-84108

End value

| Parameter Name            | Begin value | End value<br>(if different) |
|---------------------------|-------------|-----------------------------|
| -----                     | -----       | -----                       |
| _spin_count               | 2000        |                             |
| cursor_sharing            | EXACT       |                             |
| db_block_size             | 16384       |                             |
| db_cache_size             | 46170898432 |                             |
| db_keep_cache_size        | 5368709120  |                             |
| db_name                   | cc1         |                             |
| db_writer_processes       | 4           |                             |
| java_pool_size            | 167772160   |                             |
| large_pool_size           | 1073741824  |                             |
| olap_page_pool_size       | 33554432    |                             |
| open_cursors              | 900         |                             |
| pga_aggregate_target      | 5368709120  |                             |
| processes                 | 700         |                             |
| sessions                  | 1000        |                             |
| sga_max_size              | 68719476736 |                             |
| sga_target                | 68719476736 |                             |
| shared_pool_reserved_size | 262144000   |                             |
| shared_pool_size          | 4294967296  |                             |





# Quick Word on RAC

- The TCP buffers are the biggest memory issue with RAC performance, size them big
- The global dictionary takes memory out of the shared pool, the more memory in the db caches, the more is needed in the shared pools to track it.



# Questions/Comments?





Thank You!

Mike Ault

[Oracle@RamSan.com](mailto:Oracle@RamSan.com)

[www.ramsan.com](http://www.ramsan.com)

[www.statspackanalyzer.com](http://www.statspackanalyzer.com)



# Using AWR For Memory Analysis


Mike Ault, Oracle Guru  
May, 2011

The World's Fastest Storage® for over thirty years!



# Michael R. Ault, Oracle Guru

- Nuclear Navy 6 years
- Nuclear Chemist/Programmer 10 years
- Kennedy Western University Graduate
- Bachelors Degree Computer Science
- Certified in all Oracle Versions Since 6
- Oracle DBA, author, since 1990





ORACLE | CERTIFIED  
PROFESSIONAL

Texas Memory Systems, Inc.

The World's Fastest Storage®



## Books by Michael R. Ault



Texas Memory Systems, Inc.

The World's Fastest Storage®



# StatspackAnalyzer.com

## Free Statspack/AWR Analysis

Sponsored by Texas Memory Systems

- Looks for IO bottlenecks and other configuration issues.
- Straightforward tuning advice



Texas Memory Systems, Inc.

The World's Fastest Storage®



## Preparation for Analysis

- Know your systems normal performance fingerprint
- Be familiar with Concepts and Tuning Guides
- Have “normal” AWR/Statspacks for comparison





# Oracle and memory

- DB cache – other than direct read/write, all data, index and undo go through here
- Shared Pool – SQL area, PL/SQL library, dictionary caches and lots more
- Streams pool – Only if streams are used
- Java Pool – usually small
- PGA – Each process gets a PGA, stores cursor and other process related information
- Log buffers – Circular buffers for redo information



# DB Cache

- Default – should be largest area
- Recycle – for frequently scanned large objects
- Keep – For frequently accessed small objects
- 2-32K areas – Was originally for TTS, now used to tune items (usually in RAC)



# Shared Pool

- Shared SQL, PL/SQL, dictionary cache plus
- 34 other areas in 9i
- 551 in 10gR2 (non-RAC) 670 (with RAC)
- 878 in 11gR2 (non-RAC)
  - Report only shows those that change
  - There have been bugs with space leaks in shared pool sub-pools



## Top-Down Approach

- Report starts with settings overview
- Next provides Top-5 waits
- Use the Waits to guide further investigation



## WORKLOAD REPOSITORY report for

| DB Name     | DB Id             | Instance           | Inst Num | Startup Time    | Release            | RAC     |
|-------------|-------------------|--------------------|----------|-----------------|--------------------|---------|
| AULTDB      | 4030696936        | aultdb1            | 1        | 04-Aug-08 10:16 | 11.1.0.6.0         | YES     |
| Host Name   | Platform          |                    |          | CPUs Cores      | Sockets Memory(GB) |         |
| aultlinux3  | Linux             | IA (32-bit)        |          | 2               | 1                  | 2.97    |
|             | Snap Id           | Snap Time          | Sessions | Curs/Sess       |                    |         |
| Begin Snap: | 91                | 04-Aug-08 12:00:15 | 41       | 1.2             |                    |         |
| End Snap:   | 92                | 04-Aug-08 13:00:28 | 47       | 1.1             |                    |         |
| Elapsed:    |                   | 60.22 (mins)       |          |                 |                    |         |
| DB Time:    |                   | 139.52 (mins)      |          |                 |                    |         |
| Cache Sizes |                   | Begin              | End      |                 |                    |         |
|             | Buffer Cache:     | 1,312M             | 1,312M   | Std Block Size: |                    | 8K      |
|             | Shared Pool Size: | 224M               | 224M     | Log Buffer:     |                    | 10,604K |



## Signs of Memory Issues

- High sequential reads
- Excessive library latches
- Large number of sorts/ hashes/GTT/bitmap ops to disk
- Large amount of IO to the temporary tablespace
- Buffer busy waits with free buffer waits
- Indications in Cache, shared, streams or java pool advisors
- Excessive reparsing and reloads of SQL and PL/SQL
- High percentage of use in the shared pool
- High CPU cycles



## Load Profile Section

| Load Profile      | Per Second | Per Transaction | Per Exec | Per Call |
|-------------------|------------|-----------------|----------|----------|
| DB Time(s):       | 2.3        | 7.1             | 0.63     | 1.05     |
| DB CPU(s):        | 0.3        | 0.9             | 0.07     | 0.13     |
| Redo size:        | 800.5      | 2,461.8         |          |          |
| Logical reads:    | 6,307.6    | 19,396.7        |          |          |
| Block changes:    | 3.6        | 10.9            |          |          |
| Physical reads:   | 2,704.9    | 8,317.8         |          |          |
| Physical writes:  | 86.9       | 267.3           |          |          |
| User calls:       | 2.2        | 6.8             |          |          |
| Parses:           | 2.0        | 6.1             |          |          |
| Hard parses:      | 0.0        | 0.1             |          |          |
| W/A MB processed: | 932,965.4  | 2,868,990.9     |          |          |
| Logons:           | 0.1        | 0.2             |          |          |
| Executes:         | 3.7        | 11.3            |          |          |
| Rollbacks:        | 0.1        | 0.3             |          |          |
| Transactions:     | 0.3        |                 |          |          |

Texas Memory Systems, Inc.

The World's Fastest Storage®



## What Are Your Efficiencies

- Should be close to 100%
- Parse issues usually are a result of:
  - Bad bind variable usage
  - Insufficient memory
  - Will also be co-indicated by low percentage of memory for multiple SQL execution





## Load Profile Section

### Instance Efficiency Percentages (Target 100%)

|                               |        |                   |        |
|-------------------------------|--------|-------------------|--------|
| Buffer Nowait %:              | 100.00 | Redo Nowait %:    | 99.97  |
| Buffer Hit %:                 | 96.09  | In-memory Sort %: | 100.00 |
| Library Hit %:                | 98.17  | Soft Parse %:     | 97.88  |
| Execute to Parse %:           | 45.80  | Latch Hit %:      | 99.95  |
| Parse CPU to Parse Elapsed %: | 0.00   | % Non-Parse CPU:  | 99.77  |

| Shared Pool Statistics     | Begin | End   |
|----------------------------|-------|-------|
| Memory Usage %:            | 81.53 | 85.39 |
| % SQL with executions>1:   | 79.29 | 79.48 |
| % Memory for SQL w/exec>1: | 76.73 | 78.19 |



## Top 5 Waits Section

- Critical to look closely at this section
- Use highest wait times to guide investigation
  - DB FILE type waits – physical IO
  - BUFFER type waits – Logical IO
  - LOG type waits – Redo related
  - PX – Parallel Query
  - GC – Global Cache (RAC related)
  - Undo – Undo or rollback segment related



## Top 5 waits section With possible cache starvation

### Top 5 Timed Foreground Events

| Event                   | Waits   | Time(s) | Avg<br>wait<br>(ms) | % DB<br>time wait | Class    |
|-------------------------|---------|---------|---------------------|-------------------|----------|
| db file sequential read | 465,020 | 3,969   | 9                   | 47.4              | User I/O |
| DB CPU                  |         | 995     |                     | 11.9              |          |
| db file parallel read   | 2,251   | 322     | 143                 | 3.8               | User I/O |
| db file scattered read  | 15,268  | 153     | 10                  | 1.8               | User I/O |
| gc current block 2-way  | 108,739 | 116     | 1                   | 1.4               | Cluster  |



## Top 5 Waits Section With Shared Pool Issues

### Top 5 Timed Events

| Event                    | Waits       | Time (s) | Avg<br>wait<br>(ms) | %Total<br>Call<br>Time | Wait Class |
|--------------------------|-------------|----------|---------------------|------------------------|------------|
| CPU time                 |             | 435,461  |                     | 41.1                   |            |
| PX Deq Credit: send blkd | 124,829,330 | 138,223  | 1                   | 13.0                   | Other      |
| library cache pin        | 20,347      | 57,692   | 2835                | 5.4                    | Concurrenc |
| library cache lock       | 19,226      | 56,078   | 2917                | 5.3                    | Concurrenc |
| db file sequential read  | 16,798,329  | 42,215   | 3                   | 4.0                    | User I/O   |

### Top 5 Timed Events

| Event                   | Waits     | Time (s) | Avg<br>wait<br>(ms) | %Total<br>Call<br>Time | Wait Class |
|-------------------------|-----------|----------|---------------------|------------------------|------------|
| CPU time                |           | 24,956   |                     | 29.3                   |            |
| latch: library cache    | 1,757,331 | 9,886    | 6                   | 11.6                   | Concurrenc |
| db file sequential read | 759,605   | 6,146    | 8                   | 7.2                    | User I/O   |
| cursor: pin S           | 2,103,389 | 4,988    | 2                   | 5.9                    | Other      |
| log file sync           | 250,039   | 2,387    | 10                  | 2.8                    | Commit     |



## Buffer Type Waits

- latch: cache buffers chains - Hot blocks, check for hot objects, high IO rates
- free buffer waits - Insufficient buffers, processes holding buffers too long, IO subsystem over loaded
- buffer busy waits - See what is causing them further along in report
- gc buffer busy - Overloaded interconnect, find problem objects and tune
- log buffer space - High load, too small a log buffer, increase log buffer size
- latch: cache buffers lru chain - Freelist issues, hot blocks, new buffers, buffers being written
- latch: cache buffer handles - Freelist issues, hot blocks
- buffer busy - See what is causing them further along in report
- no free buffers - Insufficient buffers, dbwr contention
- Free buffer waits - insufficient buffers



# Fixing Cache Waits

- Reduce logical IO rates (buffer caches latch)
- Increase the cache size (lru chain latch)
- Increase the cache size (free buffer waits)
- Increase `_db_block_lru_latches`
- Increase `_db_block_hash_buckets`
- Reduce hot blocks



# Shared Pool Waits

- library cache pin – Loading or compiling same SQL
- library cache lock – Loading or compiling same SQL
- \*latch: library cache – Usually a result of excessive parsing
- latch: shared pool latch – Parsing issues
- \*latch: library cache lock – Usually a result of excessive parsing
- \*latch: library cache – Usually a result of excessive parsing
- row cache lock – shared pool too small
- Library cache load lock – Wait for a reload by another session. Excessive hard/soft parsing.
- \* Gone in 11g to mutex



# Shared Pool Mutexes

- Cursor:mutex X - resource is busy, requestor needs exclusive access
- Cursor:mutex S - resource is held in X mode by another session
- Cursor:pin X - resource is held in S or X by another session
- Cursor:pin S - re-execute of same cursor
- Cursor:pin S wait on X - resource is held in X mode by another session
- Library cache: mutex X - Bind variable issues
- Library cache: mutex S - Bind variable issues
- Less costly than latches





# What to Do?

- Share cursors (avoid hard parsing)
  - BIND VARIABLES!!!!!!
  - Cursor\_sharing
- Avoid soft parsing
  - Cursor\_space\_for\_time
  - Session\_cached\_cursors
- Avoid invalidations and reloads
  - Make sure shared pool is large enough



## What Next?

- Determine wait events of concern
- Drill down to specific sections of report for deeper analysis
- Use custom scripts, ADDM and Ash to investigate issues



# Classes

Wait Class

DB/Inst: Snaps: 84084-84108

-> s - second  
-> cs - centisecond - 100th of a second  
-> ms - millisecond - 1000th of a second  
-> us - microsecond - 1000000th of a second  
-> ordered by wait time desc, waits desc

| Wait Class     | Waits       | %Time<br>-outs | Total Wait<br>Time (s) | Avg<br>wait<br>(ms) | Waits<br>/txn |
|----------------|-------------|----------------|------------------------|---------------------|---------------|
| Other          | 153,619,985 | 16.5           | 192,921                | 1                   | 102.3         |
| Concurrency    | 2,536,362   | 26.9           | 128,816                | 51                  | 1.7           |
| User I/O       | 30,594,385  | .0             | 124,207                | 4                   | 20.4          |
| System I/O     | 5,104,873   | .0             | 17,633                 | 3                   | 3.4           |
| Application    | 65,645      | 5.0            | 6,508                  | 99                  | 0.0           |
| Commit         | 267,317     | .0             | 4,234                  | 16                  | 0.2           |
| Configuration  | 553,825     | 69.5           | 858                    | 2                   | 0.4           |
| Network        | 13,513,847  | .0             | 274                    | 0                   | 9.0           |
| Administrative | 30          | 70.0           | 0                      | 10                  | 0.0           |



# Operating System Statistics

Operating System Statistics

DB/Inst: Snaps: 84084-84108

| Statistic              | Total           |
|------------------------|-----------------|
| BUSY_TIME              | 45,601,415      |
| IDLE_TIME              | 6,316,939       |
| IOWAIT_TIME            | 567,343         |
| NICE_TIME              | 810,986         |
| SYS_TIME               | 3,169,946       |
| USER_TIME              | 41,265,848      |
| LOAD                   | 50              |
| RSRC_MGR_CPU_WAIT_TIME | 0               |
| PHYSICAL_MEMORY_BYTES  | 270,208,987,136 |
| NUM_CPUS               | 24              |
| NUM_CPU_SOCKETS        | 4               |

-----  
Texas Memory Systems, Inc.

-----  
The World's Fastest Storage®



# SQL Areas

**SQL ordered by CPU Time - Sorting, bad paths**

**SQL ordered by Gets - Excessive logical IO**

**SQL ordered by Reads - Cache starvation**

**SQL ordered by Parse Calls - Cursor sharing, cursor caching**

**SQL ordered by Version Count - Versioning is usually due to a bug, check with support**

- **Tune SQL that appears in more than one of these areas**
- **Tune SQL at the top of these sections**

Texas Mining Systems, Inc. is a registered provider of Oracle® Database Storage®



# System Statistics

- Many-many statistics
- Many are not useful to the DBA
- Many are useless for memory area tuning
- Many give ideas of how memory is used, but not how to tune it
- Usually these will confirm what you have already found



# System Statistics

Instance Activity Stats

DB/Inst: Snaps: 84084-84108

| Statistic                        | Total         | per Second | per Trans |
|----------------------------------|---------------|------------|-----------|
| dirty buffers inspected          | 686,267       | 31.8       | 0.5       |
| execute count                    | 78,907,090    | 3,656.2    | 52.6      |
| free buffer inspected            | 161,591,258   | 7,487.4    | 107.6     |
| free buffer requested            | 176,367,274   | 8,172.1    | 117.5     |
| hot buffers moved to head of LRU | 15,346,759    | 711.1      | 10.2      |
| immediate (CR) block cleanout ap | 2,267,512     | 105.1      | 1.5       |
| immediate (CURRENT) block cleano | 5,139,016     | 238.1      | 3.4       |
| no buffer to keep pinned count   | 136,849       | 6.3        | 0.1       |
| no work - consistent read gets   | 4,459,140,613 | 206,616.1  | 2,969.8   |
| opened cursors cumulative        | 26,795,933    | 1,241.6    | 17.9      |
| parse count (failures)           | 160           | 0.0        | 0.0       |
| parse count (hard)               | 398,147       | 18.5       | 0.3       |
| parse count (total)              | 20,200,501    | 936.0      | 13.5      |
| parse time cpu                   | 3,883,178     | 179.9      | 2.6       |
| parse time elapsed               | 7,474,786     | 346.4      | 5.0       |
| physical read total IO requests  | 37,303,810    | 1,728.5    | 24.8      |
| physical reads direct temporary  | 58,378,313    | 2,705.0    | 38.9      |
| physical write total IO requests | 13,738,098    | 636.6      | 9.2       |
| physical writes direct temporary | 58,440,795    | 2,707.9    | 38.9      |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# System Statistics

|                                  |                 |             |          |
|----------------------------------|-----------------|-------------|----------|
| pinned buffers inspected         | 120,843         | 5.6         | 0.1      |
| recursive calls                  | 749,184,714     | 34,713.8    | 499.0    |
| recursive cpu usage              | 39,323,240      | 1,822.1     | 26.2     |
| redo log space requests          | 190             | 0.0         | 0.0      |
| redo log space wait time         | 333             | 0.0         | 0.0      |
| redo synch time                  | 433,625         | 20.1        | 0.3      |
| redo synch writes                | 236,148         | 10.9        | 0.2      |
| redo write time                  | 567,670         | 26.3        | 0.4      |
| redo writer latching time        | 56,827          | 2.6         | 0.0      |
| redo writes                      | 1,127,300       | 52.2        | 0.8      |
| rollback changes - undo records  | 1,395,329       | 64.7        | 0.9      |
| rollbacks only - consistent read | 346,504         | 16.1        | 0.2      |
| session cursor cache hits        | 21,520,355      | 997.2       | 14.3     |
| session logical reads            | 10,474,545,504  | 485,342.3   | 6,976.0  |
| sorts (disk)                     | 3,529           | 0.2         | 0.0      |
| sorts (memory)                   | 9,012,270       | 417.6       | 6.0      |
| sorts (rows)                     | 110,063,794,220 | 5,099,850.2 | 73,302.3 |
| sql area evicted                 | 327,084         | 15.2        | 0.2      |
| sql area purged                  | 29,720          | 1.4         | 0.0      |
| table scans (long tables)        | 1,149,945       | 53.3        | 0.8      |
| table scans (short tables)       | 7,528,140       | 348.8       | 5.0      |
| transaction rollbacks            | 252,407         | 11.7        | 0.2      |

Texas Memory Systems, Inc.

The World's Fastest Storage®





# System Statistics

|                                 |            |       |     |
|---------------------------------|------------|-------|-----|
| user I/O wait time              | 12,422,069 | 575.6 | 8.3 |
| user calls                      | 8,038,839  | 372.5 | 5.4 |
| user commits                    | 1,439,821  | 66.7  | 1.0 |
| user rollbacks                  | 61,684     | 2.9   | 0.0 |
| workarea executions - multipass | 0          | 0.0   | 0.0 |
| workarea executions - onepass   | 5,293      | 0.3   | 0.0 |
| workarea executions - optimal   | 7,113,060  | 329.6 | 4.7 |



# Instance Activity Statistics

Instance Activity Stats - Absolute

-> Statistics with absolute values (should not be diffed)

| Statistic                  | Begin Value | End Value  |
|----------------------------|-------------|------------|
| -----                      | -----       | -----      |
| session cursor cache count | 28,024,069  | 28,789,659 |
| opened cursors current     | 2,921       | 6,982      |
| workarea memory allocated  | 289,532     | 2,531,741  |
| logons current             | 144         | 287        |
| -----                      | -----       | -----      |



# Tablespace/File IO Reports

- Helps confirm IO issues
- Also helps with temp area issue determination



# Tablespace IO

Tablespace IO Stats  
-> ordered by IOs (Reads + Writes) desc

Tablespace

|                         | Av<br>Reads | Av<br>Reads/s | Av<br>Rd(ms) | Av<br>Blks/Rd | Writes    | Av<br>Writes/s | Buffer<br>Waits | Av<br>Buf<br>Wt(ms) |
|-------------------------|-------------|---------------|--------------|---------------|-----------|----------------|-----------------|---------------------|
| TEMP                    | 11,484,000  | 532           | 16.3         | 4.1           | 3,478,365 | 161            | 12,266          | 2.0                 |
| REPMAN_TEMP             | 1,703,767   | 79            | 27.2         | 8.2           | 1,457,241 | 68             | 0               | 0.0                 |
| UNDOTBS3                | 30,012      | 1             | 8.0          | 1.0           | 1,512,571 | 70             | 142,889         | 1.1                 |
| RSNET_DTSA              | 1,496,441   | 69            | 1.2          | 2.0           | 2,454     | 0              | 130,753         | 1.3                 |
| LOREAL_D_CVS_DAILY_ITSA | 846,665     | 39            | 0.9          | 1.0           | 338       | 0              | 0               | 0.0                 |



# Buffer Pool Statistics

Buffer Pool Statistics

DB/Inst: CC1/cc1 Snaps: 84084-84108

-> Standard block size Pools D: default, K: keep, R: recycle

-> Default Pools for other block sizes: 2k, 4k, 8k, 16k, 32k

| P | Number of Pool<br>Buffers | Hit% | Buffer<br>Gets | Physical<br>Reads | Physical<br>Writes | Free Buff<br>Wait | Write Comp<br>Wait | Buffer<br>Busy<br>Waits |
|---|---------------------------|------|----------------|-------------------|--------------------|-------------------|--------------------|-------------------------|
| D | 3,361,107                 | 96   | 3,643,978,600  | 163,055,679       | 14,338,623         | 0                 | 0                  | 711,281                 |
| K | 321,600                   | 100  | 2,527,600,634  | 7,379             | 28,755             | 0                 | 0                  | 123                     |

- Note that there are Buffer Busy Waits, but no Free Buffer Waits
- These are due to hot block contention
- Increasing memory probably won't help with this
- However...also look at db file sequential read waits and the cache advisory section



# Buffer Pool Advisory Section

Buffer Pool Advisory

-> Only rows with estimated physical reads >0 are displayed

-> ordered by Block Size, Buffers For Estimate

| P   | Size for<br>Est (M) | Size<br>Factor | Buffers for<br>Estimate | Est<br>Phys<br>Read<br>Factor | Estimated<br>Physical Reads |
|-----|---------------------|----------------|-------------------------|-------------------------------|-----------------------------|
| D   | 5,344               | .1             | 335,670                 | 1.9                           | 15,767,325,073              |
| D   | 10,688              | .2             | 671,340                 | 1.4                           | 11,371,357,960              |
| ... |                     |                |                         |                               |                             |
| D   | 106,880             | 2.0            | 6,713,400               | 1.0                           | 7,964,367,701               |
| K   | 512                 | .1             | 32,160                  | 102.8                         | 3,507,100,178               |
| K   | 1,024               | .2             | 64,320                  | 7.8                           | 264,615,629                 |
| K   | 1,536               | .3             | 96,480                  | 1.4                           | 49,384,590                  |
| ... |                     |                |                         |                               |                             |
| K   | 10,240              | 2.0            | 643,200                 | 1.0                           | 32,639,643                  |

Texas Memory Systems, Inc.

The World's Fastest Storage®



## Buffer Pool Advisory Section

- As you can see, this report shows even doubling the default or keep would be no benefit
- Let's look at one that would benefit from increased buffer pool



# Buffer Pool Advisor Section

## Buffer Pool Statistics

DB/Inst: Snaps: 26064-26097

-> Standard block size Pools D: default, K: keep, R: recycle

-> Default Pools for other block sizes: 2k, 4k, 8k, 16k, 32k

| P | Number of Pool<br>Buffers Hit% | Buffer<br>Gets | Physical<br>Reads | Physical<br>Writes | Free Buff<br>Wait | Writ<br>Comp<br>Wait | Buffer<br>Busy<br>Waits |
|---|--------------------------------|----------------|-------------------|--------------------|-------------------|----------------------|-------------------------|
| D | 818,201 99                     | 130,795,544    | 1,578,580         | 276,075            | 0                 | 0                    | 3,418                   |

- Before we go there...look here
- Notice no free buffer waits





## Buffer Pool Advisor Section

Buffer Pool Advisory

Snap: 26097

-> Only rows with estimated physical reads >0 are displayed

-> ordered by Block Size, Buffers For Estimate

| P   | Size for<br>Est (M) | Size<br>Factor | Buffers for<br>Estimate | Est<br>Phys<br>Read<br>Factor | Estimated<br>Physical Reads |
|-----|---------------------|----------------|-------------------------|-------------------------------|-----------------------------|
| D   | 656                 | .1             | 81,139                  | 2.0                           | 125,592,784                 |
| D   | 1,312               | .2             | 162,278                 | 1.8                           | 113,080,052                 |
| ... |                     |                |                         |                               |                             |
| D   | 10,496              | 1.6            | 1,298,224               | 0.7                           | 42,942,366                  |
| D   | 11,152              | 1.7            | 1,379,363               | 0.7                           | 41,649,501                  |
| D   | 11,808              | 1.8            | 1,460,502               | 0.6                           | 40,403,058                  |

- Notice that at 1.8 times the current size physical reads down by 40%

- ~~Do file sequential reads was 13% of waits~~

The World's Fastest Storage®



# PGA Analysis

- Several of the next sections deal with PGA
- PGA\_AGGREGATE\_TARGET sets the PGA area
- 5% of PGA\_AGGREGATE\_TARGET can be allocated to each session up to a maximum of “\_PGA\_MAX\_SIZE” which is usually 200 or 500 megabytes
- Manually setting SORT\_AREA\_SIZE or HASH\_AREA\_SIZE overrides at the session level
- Some processes such as RMAN and shared servers don't use PGA\_AGGREGATE\_TARGET but use the old manual settings, indicated by 4-8 or 8-16m sorts even with adequate PGA\_AGGREGATE\_TARGET settings



# PGA Analysis

PGA Aggr Summary

DB/Inst: Snaps: 84084-84108

-> PGA cache hit % - percentage of W/A (WorkArea) data processed only in-memory

| PGA Cache Hit % | W/A MB Processed | Extra W/A MB Read/Written |
|-----------------|------------------|---------------------------|
| 82.1            | 4,495,435        | 979,073                   |



# PGA Analysis

PGA Aggr Target Stats

DB/Inst: Snaps: 84084-84108

- > B: Begin snap E: End snap (rows identified with B or E contain data which is absolute i.e. not diffed over the interval)
- > Auto PGA Target - actual workarea memory target
- > W/A PGA Used - amount of memory used for all Workareas (manual + auto)
- > %PGA W/A Mem - percentage of PGA memory allocated to workareas
- > %Auto W/A Mem - percentage of workarea memory controlled by Auto Mem Mgmt
- > %Man W/A Mem - percentage of workarea memory under manual control

|   | PGA Aggr<br>Target(M) | Auto PGA<br>Target(M) | PGA Mem<br>Alloc(M) | W/A PGA<br>Used(M) | %PGA<br>W/A<br>Mem | %Auto<br>W/A<br>Mem | %Man<br>W/A<br>Mem | Global Mem<br>Bound(K) |
|---|-----------------------|-----------------------|---------------------|--------------------|--------------------|---------------------|--------------------|------------------------|
| B | 5,120                 | 4,320                 | 1,680.5             | 193.5              | 11.5               | 99.7                | .3                 | 524,280                |
| E | 5,120                 | 4,202                 | 4,400.5             | 2,219.2            | 50.4               | 99.9                | .1                 | 524,280                |



# PGA Analysis

PGA Aggr Target Histogram

Snapshots: 84084-84108

Low High

| Optimal | Optimal | Total Execs | Optimal Execs | 1-Pass Execs | M-Pass Execs |
|---------|---------|-------------|---------------|--------------|--------------|
| 2K      | 4K      | 6,308,435   | 6,308,435     | 0            | 0            |
| 64K     | 128K    | 32,500      | 32,500        | 0            | 0            |
| 128K    | 256K    | 36,535      | 36,535        | 0            | 0            |
| 256K    | 512K    | 47,477      | 47,477        | 0            | 0            |
| 512K    | 1024K   | 353,344     | 353,344       | 0            | 0            |
| 1M      | 2M      | 201,558     | 201,558       | 0            | 0            |
| 2M      | 4M      | 22,468      | 22,468        | 0            | 0            |
| 4M      | 8M      | 21,796      | 21,725        | 71           | 0            |
| 8M      | 16M     | 28,892      | 28,714        | 178          | 0            |
| 16M     | 32M     | 30,478      | 30,346        | 132          | 0            |
| 32M     | 64M     | 19,898      | 18,690        | 1,208        | 0            |
| 64M     | 128M    | 9,080       | 7,284         | 1,796        | 0            |
| 128M    | 256M    | 1,682       | 732           | 950          | 0            |
| 256M    | 512M    | 734         | 179           | 555          | 0            |
| 512M    | 1024M   | 329         | 58            | 271          | 0            |
| 1G      | 2G      | 131         | 14            | 117          | 0            |
| 2G      | 4G      | 17          | 4             | 13           | 0            |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# PGA Analysis

- $SS * 100 / 2.5 * NOS$ 
  - SS- sort size
  - NOS - number of sorts
- $128M * 100 / 2.5 = 5120$
- Hash gets full 5%
- Sort gets  $\frac{1}{2}$  of hash
- $\frac{1}{2}$  of 5 is 2.5
- 10|10 rule:
  - 10% of sessions are active
  - 10% of active sessions doing sorts
- 10% of 208 sessions is 20.8; 10% of 20.8 is 2
- $2 * 5120$  is 10420 gb
- Let's see what the advisor says...



# PGA Analysis

PGA Memory Advisory

Snap: 84108

-> When using Auto Memory Mgmt, minimally choose a pga\_aggregate\_target value where Estd PGA Overalloc Count is 0

| PGA Target<br>Est (MB) | Size<br>Factr | W/A MB<br>Processed | Estd Extra<br>W/A MB Read/<br>Written to Disk | Estd PGA<br>Cache<br>Hit % | Estd PGA<br>Overalloc<br>Count |
|------------------------|---------------|---------------------|-----------------------------------------------|----------------------------|--------------------------------|
| 640                    | 0.1           | 167,211,258.2       | 133,135,735.1                                 | 56.0                       | 18,399                         |
| 1,280                  | 0.3           | 167,211,258.2       | 65,323,839.1                                  | 72.0                       | 1,933                          |
| 2,560                  | 0.5           | 167,211,258.2       | 34,248,434.3                                  | 83.0                       | 1,842                          |
| 3,840                  | 0.8           | 167,211,258.2       | 28,768,442.5                                  | 85.0                       | 1,803                          |
| 5,120                  | 1.0           | 167,211,258.2       | 19,597,963.1                                  | 90.0                       | 860                            |
| 6,144                  | 1.2           | 167,211,258.2       | 14,425,059.9                                  | 92.0                       | 845                            |
| 7,168                  | 1.4           | 167,211,258.2       | 13,624,988.5                                  | 92.0                       | 839                            |
| 8,192                  | 1.6           | 167,211,258.2       | 12,882,080.7                                  | 93.0                       | 0                              |
| 9,216                  | 1.8           | 167,211,258.2       | 12,146,846.2                                  | 93.0                       | 0                              |
| 10,240                 | 2.0           | 167,211,258.2       | 11,689,119.6                                  | 93.0                       | 0                              |
| 15,360                 | 3.0           | 167,211,258.2       | 10,710,132.3                                  | 94.0                       | 0                              |
| 20,480                 | 4.0           | 167,211,258.2       | 10,563,396.2                                  | 94.0                       | 0                              |
| 30,720                 | 6.0           | 167,211,258.2       | 10,539,291.8                                  | 94.0                       | 0                              |
| 40,960                 | 8.0           | 167,211,258.2       | 10,539,291.8                                  | 94.0                       | 0                              |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# PGA Analysis

- Advisor says 8 gb
- Calculation says 10 gb
- Take whichever you feel will give best results
- Won't be used unless it needs to be





## Shared Pool Advisor

- Use reloads and used percentages to guide you
- The advisor rarely has any meaningful information
- Even when reloads are huge and other factors show the pool should be increased, or, decreased it has told me things where fine



# Shared Pool Advisor

Shared Pool Advisory

DB/Inst: Snap: 84108

-> SP: Shared Pool Est LC: Estimated Library Cache Factr: Factor

-> Note there is often a 1:Many correlation between a single logical object in the Library Cache, and the physical number of memory objects associated with it. Therefore comparing the number of Lib Cache objects (e.g. in v\$librarycache), with the number of Lib Cache Memory Objects is invalid.

| Shared<br>Pool<br>Size(M) | SP<br>Size<br>Factr | Est LC<br>Size<br>(M) | Est LC<br>Mem Obj | Est LC<br>Time<br>Saved<br>(s) | Est LC<br>Time<br>Saved<br>Factr | Est LC<br>Load<br>Time<br>(s) | Est LC<br>Load<br>Time<br>Factr | Est LC<br>Mem<br>Obj Hits |
|---------------------------|---------------------|-----------------------|-------------------|--------------------------------|----------------------------------|-------------------------------|---------------------------------|---------------------------|
| 2,160                     | .4                  | 619                   | 76,837            | #####                          | 1.0                              | #####                         | 2.4                             | 88,538,740                |
| 2,736                     | .5                  | 1,188                 | 95,749            | #####                          | 1.0                              | #####                         | 2.1                             | 88,936,333                |
| 3,312                     | .6                  | 1,761                 | 110,785           | #####                          | 1.0                              | #####                         | 1.8                             | 89,297,339                |
| 3,888                     | .7                  | 2,333                 | 125,755           | #####                          | 1.0                              | #####                         | 1.6                             | 89,610,155                |
| ...                       |                     |                       |                   |                                |                                  |                               |                                 |                           |
| 8,496                     | 1.5                 | 6,916                 | 238,592           | #####                          | 1.0                              | #####                         | .5                              | 91,076,008                |
| 9,072                     | 1.6                 | 7,489                 | 252,248           | #####                          | 1.0                              | #####                         | .4                              | 91,187,541                |
| 9,648                     | 1.7                 | 8,061                 | 264,748           | #####                          | 1.0                              | #####                         | .3                              | 91,291,114                |
| 10,224                    | 1.8                 | 8,632                 | 274,837           | #####                          | 1.0                              | #####                         | .3                              | 91,388,340                |
| 10,800                    | 1.9                 | 9,201                 | 284,723           | #####                          | 1.0                              | #####                         | .2                              | 91,480,577                |
| 11,376                    | 2.0                 | 9,774                 | 293,073           | #####                          | 1.0                              | #####                         | .2                              | 91,569,191                |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# SGA Target Advisor

- I believe you should set the base parameters and then allow SGA\_TARGET to set itself
- Other calculation schemes involve adding up the proposed sizes needed and then setting the value
- Same logic applies to MEMORY\_TARGET only include PGA\_AGGREGATE\_TARGET in the base
- Then set SGA\_MAX\_SIZE or MEMORY\_MAX\_SIZE 10-20% higher than their associated TARGET values.



# SGA Target

SGA Target Advisory

Snap: 26097

| SGA Target<br>Size (M) | SGA Size<br>Factor | Est DB<br>Time (s) | Est Physical<br>Reads |
|------------------------|--------------------|--------------------|-----------------------|
| 2,048                  | 0.3                | 8,957,297          | 125,595,242           |
| 4,096                  | 0.5                | 745,014            | 99,011,933            |
| 6,144                  | 0.8                | 683,241            | 84,201,232            |
| 8,192                  | 1.0                | 603,250            | 62,255,994            |
| 10,240                 | 1.3                | 573,691            | 48,902,083            |
| 12,288                 | 1.5                | 573,630            | 42,944,185            |
| 14,336                 | 1.8                | 573,630            | 37,745,809            |
| 16,384                 | 2.0                | 573,630            | 37,745,809            |

- At 2x shows a 60% reduction in PR, same as doubling the cache for this instance



# Streams Pool Advisor

- Only valid if you use streams
- If it shows you have spills to disk  
make pool larger



# Streams Pool Advisor

Streams Pool Advisory

DB/Inst: Snap: 84108

| Size for<br>Est (MB) | Size<br>Factor | Est Spill<br>Count | Est Spill<br>Time (s) | Est Unspill<br>Count | Est Unspill<br>Time (s) |
|----------------------|----------------|--------------------|-----------------------|----------------------|-------------------------|
| 16                   | 0.5            | 0                  | 0                     | 0                    | 0                       |
| 32                   | 1.0            | 0                  | 0                     | 0                    | 0                       |
| 48                   | 1.5            | 0                  | 0                     | 0                    | 0                       |
| 64                   | 2.0            | 0                  | 0                     | 0                    | 0                       |
| 80                   | 2.5            | 0                  | 0                     | 0                    | 0                       |
| ...                  |                |                    |                       |                      |                         |
| 240                  | 7.5            | 0                  | 0                     | 0                    | 0                       |
| 256                  | 8.0            | 0                  | 0                     | 0                    | 0                       |
| 272                  | 8.5            | 0                  | 0                     | 0                    | 0                       |
| 288                  | 9.0            | 0                  | 0                     | 0                    | 0                       |
| 304                  | 9.5            | 0                  | 0                     | 0                    | 0                       |
| 320                  | 10.0           | 0                  | 0                     | 0                    | 0                       |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Java Pool Advisor

- Another area usually not used
- If you use it you will get errors if it is too small
- Under AMM will grow but usually won't shrink



# Java Pool Advisor

Java Pool Advisory

DB/Inst: Snap: 37

| Java<br>Pool<br>Size(M) | JP<br>Size<br>Factr | Est LC<br>Size<br>(M) | Est LC<br>Mem Obj | Est LC<br>Time<br>Saved<br>(s) | Est LC<br>Time<br>Saved<br>Factr | Est LC<br>Load<br>Time<br>(s) | Est LC<br>Load<br>Time<br>Factr | Est LC<br>Mem<br>Obj Hits |
|-------------------------|---------------------|-----------------------|-------------------|--------------------------------|----------------------------------|-------------------------------|---------------------------------|---------------------------|
| 64                      | .5                  | 10                    | 168               | 10                             | 1.0                              | 11,974                        | 1.0                             | 389                       |
| 128                     | 1.0                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |
| 192                     | 1.5                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |
| 256                     | 2.0                 | 12                    | 201               | 10                             | 1.0                              | 11,974                        | 1.0                             | 465                       |





# Buffer Wait Analysis

- Look here to see what is causing buffer waits
- Data blocks usually predominate
- High buffer waits for data blocks plus high db file sequential reads will usually indicate memory starvation even without free buffer waits



# Buffer Wait Analysis

Buffer Wait Statistics

DB/Inst:

Snap: 84084-84108

-> ordered by wait time desc, waits desc

| Class             | Waits   | Total Wait Time (s) | Avg Time (ms) |
|-------------------|---------|---------------------|---------------|
| data block        | 539,313 | 1,770               | 3             |
| 1st level bmb     | 5,873   | 88                  | 15            |
| undo block        | 49,801  | 71                  | 1             |
| undo header       | 35,848  | 66                  | 2             |
| file header block | 57,833  | 39                  | 1             |
| segment header    | 12,980  | 28                  | 2             |
| 3rd level bmb     | 4,900   | 22                  | 5             |
| 2nd level bmb     | 5,206   | 20                  | 4             |
| extent map        | 73      | 0                   | 5             |



# Enqueues and Latches

- Enqueues are usually for physical objects or transactions and rarely memory related
- Latches can be used to verify findings



# Latches

## Latch Activity

DB/Inst: Snaps: 84084-84108

-> "Get Requests", "Pct Get Miss" and "Avg Slps/Miss" are statistics for willing-to-wait latch get requests

-> "NoWait Requests", "Pct NoWait Miss" are for no-wait latch get requests

-> "Pct Misses" for both should be very close to 0.0

| Latch Name               | Get Requests  | Pct      | Avg        | Wait     | NoWait Requests | Pct         |
|--------------------------|---------------|----------|------------|----------|-----------------|-------------|
|                          |               | Get Miss | Slps /Miss | Time (s) |                 | NoWait Miss |
| SGA IO buffer pool latch | 154,979       | 0.0      | 0.7        | 0        | 261,675         | 0.3         |
| SQL memory manager latch | 21,353        | 12.7     | 1.2        | 32       | 7,066           | 0.1         |
| active service list      | 2,524,503     | 1.7      | 0.2        | 10       | 7,615           | 0.1         |
| cache buffers lru chain  | 66,153,522    | 1.0      | 0.6        | 1172     | 469,078,119     | 1.8         |
| cache table scan latch   | 3,749,595     | 0.3      | 0.3        | 9        | 3,986,131       | 0.3         |
| library cache            | 272,976,941   | 0.1      | 0.4        | 2242     | 1,582,568       | 216.3       |
| library cache lock       | 179,441,868   | 0.1      | 0.1        | 34       | 58,383          | 0.1         |
| object queue header oper | 883,815,615   | 0.1      | 0.1        | 365      | 45,167          | 0.2         |
| process queue reference  | 8,224,189,548 | 0.0      | 0.0        | 15       | 613,910,934     | 4.3         |
| redo allocation          | 8,482,485     | 8.1      | 0.1        | 271      | 435,767,620     | 1.4         |
| row cache objects        | 871,266,120   | 0.4      | 0.2        | 1765     | 580,035         | 1.6         |
| simulator lru latch      | 7,890         | 1.1      | 1.1        | 1        | #####           | 17.1        |
| undo global data         | 68,788,574    | 0.0      | 0.1        | 8        | 42,419          | 0.1         |



# Latches

- When a latch can't be obtained a process will "sleep"
- The number of CPU cycles is set by the "\_SPIN\_COUNT" setting
- Defaults to 2000, was set in the era of slow CPUs
- Experts agree setting as high as 8000-10000 is acceptable
- Reduces sleeps and improved performance



# Latch Sleeps

Latch Sleep Breakdown  
-> ordered by misses desc

DB/Inst: Snaps: 84084-84108

Latch Name

| Get Requests                  | Misses     | Sleeps  | Spin Gets | Sleep1 | Sleep2 | Sleep3 |
|-------------------------------|------------|---------|-----------|--------|--------|--------|
| cache buffers chains          |            |         |           |        |        |        |
| 9,561,948                     | #####      | 402,582 | #####     | 0      | 0      | 0      |
| session allocation            |            |         |           |        |        |        |
| 570,516,196                   | 14,728,230 | 461,868 | #####     | 0      | 0      | 0      |
| messages                      |            |         |           |        |        |        |
| 105,586,860                   | 3,521,238  | 16,783  | 3,506,276 | 0      | 0      | 0      |
| row cache objects             |            |         |           |        |        |        |
| 871,266,120                   | 3,508,138  | 655,697 | 2,914,340 | 0      | 0      | 0      |
| process queue reference       |            |         |           |        |        |        |
| 8,224,189,548                 | 2,645,262  | 19,802  | 2,629,967 | 0      | 0      | 0      |
| active checkpoint queue latch |            |         |           |        |        |        |
| 50,214,009                    | 1,082,455  | 13,133  | 1,070,658 | 0      | 0      | 0      |
| redo writing                  |            |         |           |        |        |        |
| 52,201,162                    | 1,026,204  | 6,893   | 1,020,132 | 0      | 0      | 0      |
| object queue header operation |            |         |           |        |        |        |
| 883,815,615                   | 936,882    | 137,869 | 812,882   | 0      | 0      | 0      |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Latch Miss Sources

Latch Miss Sources

DB/Inst: Snaps: 84084-84108

-> only latches with sleeps are shown

-> ordered by name, sleeps desc

| Latch Name               | Where                       | Nowait<br>Misses | Sleeps  | Waiter<br>Sleeps |
|--------------------------|-----------------------------|------------------|---------|------------------|
| In memory undo latch     | ktiFlush: child             | 0                | 12,727  | 3,775            |
| ...                      |                             |                  |         |                  |
| cache buffers chains     | kcbgtcr: kslbegin excl      | 0                | 821,219 | 689,229          |
| ...                      |                             |                  |         |                  |
| cache buffers lru chain  | kcbxsv: move to being wri   | 0                | 19,728  | 87,892           |
| ...                      |                             |                  |         |                  |
| library cache            | kglpnd1: child: after proc  | 0                | 28,574  | 10,336           |
| ...                      |                             |                  |         |                  |
| library cache lock       | kgllkd1: child: cleanup     | 0                | 7,106   | 6,749            |
| ...                      |                             |                  |         |                  |
| object queue header oper | kcbw_unlink_q               | 0                | 139,971 | 27,436           |
| ...                      |                             |                  |         |                  |
| parameter table allocati | ksp_param_table_free        | 0                | 60,653  | 60,605           |
| row cache objects        | kqrpre: find obj            | 0                | 381,022 | 379,948          |
| ...                      |                             |                  |         |                  |
| session allocation       | ksuprc                      | 0                | 191,233 | 175,384          |
| ...                      |                             |                  |         |                  |
| shared pool              | kg halo                     | 0                | 194,517 | 123,739          |
| ...                      |                             |                  |         |                  |
| simulator hash latch     | <u>kcbisacc: lookup dba</u> | 0                | 16,413  | 26,348           |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Dictionary Cache

- Misses can be costly
- Only correction is larger shared pool
- Used to be individually controlled by DC parameters (v6)
- Now controlled automatically





# Dictionary Cache

Dictionary Cache Stats

DB/Inst: Snaps: 84084-84108

-> "Pct Misses" should be very low (< 2% in most cases)

-> "Final Usage" is the number of cache entries being used

| Cache                | Get<br>Requests | Pct<br>Miss | Scan<br>Reqs | Pct<br>Miss | Mod<br>Reqs | Final<br>Usage |
|----------------------|-----------------|-------------|--------------|-------------|-------------|----------------|
| dc_awr_control       | 442             | 9.3         | 0            | N/A         | 48          | 1              |
| dc_constraints       | 60,034          | 41.5        | 0            | N/A         | 60,034      | 49             |
| dc_files             | 559             | 97.3        | 0            | N/A         | 0           | 0              |
| dc_global_oids       | 20,580          | 3.0         | 0            | N/A         | 0           | 25             |
| dc_histogram_data    | 7,832,289       | 2.0         | 0            | N/A         | 19,621      | 7,472          |
| dc_histogram_defs    | 5,795,619       | 15.5        | 0            | N/A         | 322,942     | 6,953          |
| dc_object_grants     | 99,052          | 11.6        | 0            | N/A         | 0           | 370            |
| dc_object_ids        | 105,285,794     | 2.3         | 0            | N/A         | 38,220      | 66,669         |
| dc_objects           | 53,343,945      | 4.8         | 0            | N/A         | 161,568     | 8,250          |
| dc_profiles          | 60,308          | 0.1         | 0            | N/A         | 0           | 5              |
| dc_segments          | 19,064,933      | 58.4        | 0            | N/A         | 256,828     | 16,244         |
| dc_sequences         | 137,407         | 0.9         | 0            | N/A         | 137,407     | 35             |
| dc_table_scns        | 180,866         | 4.2         | 0            | N/A         | 88          | 14             |
| dc_tablespace_quotas | 101,823         | 2.5         | 0            | N/A         | 101,679     | 43             |
| dc tablespaces       | 11,521,250      | 1.3         | 0            | N/A         | 18          | 10,945         |
| dc usernames         | 750,849         | 0.4         | 0            | N/A         | 0           | 104            |
| global database name | 118             | 6.8         | 0            | N/A         | 0           | 1              |
| kqlsubheap_object    | 1               | 100.0       | 0            | N/A         | 0           | 0              |
| outstanding_alerts   | 394,019         | 41.7        | 0            | N/A         | 0           | 10,943         |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Library Cache Activity

- Usually only see issues in SQL and PL/SQL areas
- Invalidations high means DDL activity
- Reloads high means SQL or PL/SQL issues



# Library Cache Activity

Library Cache Activity

DB/Inst: Snaps: 84084-84108

-> "Pct Misses" should be very low

| Namespace       | Get<br>Requests | Pct<br>Miss | Pin<br>Requests | Pct<br>Miss | Reloads   | Invali-<br>dations |
|-----------------|-----------------|-------------|-----------------|-------------|-----------|--------------------|
| BODY            | 106,197         | 0.3         | 2,372,306       | 0.1         | 1,288     | 0                  |
| CLUSTER         | 11,232          | 0.4         | 19,082          | 0.7         | 81        | 0                  |
| INDEX           | 1,728,015       | 1.9         | 28,528,015      | 0.2         | 17,973    | 0                  |
| PIPE            | 9               | 0.0         | 9               | 0.0         | 0         | 0                  |
| SQL AREA        | 3,482,631       | 14.7        | 108,055,227     | 2.4         | 1,276,544 | #####              |
| TABLE/PROCEDURE | 290,691         | 6.1         | 32,986,432      | 0.5         | 76,695    | 0                  |
| TRIGGER         | 621,821         | 0.1         | 2,312,915       | 0.1         | 2,423     | 0                  |



# Process Statistics

## Process Memory Summary

DB/Inst: Snaps: 84084-84108

-> B: Begin snap E: End snap

-> All rows below contain absolute values (i.e. not diffed over the interval)

-> Max Alloc is Maximum PGA Allocation size at snapshot time

-> Hist Max Alloc is the Historical Max Allocation for still-connected processes

-> ordered by Begin/End snapshot, Alloc (MB) desc

| Category   | Alloc<br>(MB) | Used<br>(MB) | Avg<br>Alloc<br>(MB) | Std Dev<br>Alloc<br>(MB) | Max<br>Alloc<br>(MB) | Hist<br>Max<br>Alloc<br>(MB) | Num<br>Proc | Num<br>Alloc |
|------------|---------------|--------------|----------------------|--------------------------|----------------------|------------------------------|-------------|--------------|
| B Freeable | 981.6         | .0           | 6.2                  | 23.0                     | 141                  | N/A                          | 158         | 158          |
| Other      | 647.4         | N/A          | 3.3                  | 5.8                      | 30                   | 102                          | 195         | 191          |
| PL/SQL     | 38.8          | 7.0          | .2                   | .6                       | 3                    | 3                            | 193         | 154          |
| SQL        | 15.7          | 10.9         | .1                   | .2                       | 2                    | 1,289                        | 174         | 128          |
| E SQL      | 1,953.6       | 1,939.0      | 6.7                  | 36.8                     | 488                  | 1,289                        | 291         | 264          |
| Other      | 1,772.1       | N/A          | 5.6                  | 16.7                     | 177                  | 498                          | 315         | 311          |
| Freeable   | 640.8         | .0           | 4.1                  | 14.2                     | 92                   | N/A                          | 156         | 156          |
| PL/SQL     | 77.0          | 13.8         | .2                   | .5                       | 3                    | 13                           | 313         | 313          |



# SGA Memory Summary

SGA Memory Summary

DB/Inst: Snaps: 84084-84108

| SGA regions      | Begin Size (Bytes) | End Size (Bytes)<br>(if different) |
|------------------|--------------------|------------------------------------|
| Database Buffers | 61,035,511,808     | 61,538,828,288                     |
| Fixed Size       | 2,212,832          |                                    |
| Redo Buffers     | 14,561,280         |                                    |
| Variable Size    | 7,667,190,816      | 7,163,874,336                      |
| sum              | 68,719,476,736     |                                    |



# SGA Breakdown

SGA breakdown difference

DB/Inst: Snaps: 84084-84108

-> ordered by Pool, Name

-> N/A value for Begin MB or End MB indicates the size of that Pool/Name was insignificant, or zero in that snapshot

| Pool   | Name                      | Begin MB | End MB   | % Diff  |
|--------|---------------------------|----------|----------|---------|
| java   | free memory               | 160.0    | 160.0    | 0.00    |
| large  | free memory               | 1,022.6  | 1,020.6  | -0.19   |
| shared | CCursor                   | 483.3    | 425.2    | -12.04  |
| shared | Checkpoint queue          | 175.8    | 175.8    | 0.00    |
| shared | KGH: NO ACCESS            | 2,445.0  | 3,079.3  | 25.94   |
| shared | KQR L PO                  | 355.6    | N/A      | -100.00 |
| shared | KQR M PO                  | 430.4    | N/A      | -100.00 |
| shared | PCursor                   | 221.9    | 190.5    | -14.16  |
| shared | db_block_hash_buckets     | 180.0    | 180.0    | 0.00    |
| shared | free memory               | 2,191.3  | 2,890.1  | 31.89   |
| shared | kzctxgjsi ksuseclid memor | 1,644.5  | 1,676.7  | 1.96    |
| shared | library cache             | 334.2    | 308.2    | -7.79   |
| shared | obj stat memo             | 466.8    | 474.8    | 1.71    |
| shared | partitioning d            | 131.8    | N/A      | -100.00 |
| shared | sql area                  | 483.6    | 236.8    | -51.04  |
| stream | free memory               | 31.9     | 31.9     | 0.00    |
|        | buffer_cache              | 58,208.0 | 58,688.0 | 0.82    |
|        | fixed_sga                 | 2.1      | 2.1      | 0.00    |
|        | log_buffer                | 13.9     | 13.9     | 0.00    |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Resize Operations

- Using AMM resize operations are automatic
- If there is room to grow, shouldn't have to borrow (steal) from other areas
- If `SGA_TARGET=SGA_MAX_SIZE` or `MEMORY_TARGET=MEMORY_MAX_SIZE` then no room to grow



# Resize Operations

Memory Dynamic Components

DB/Inst: Snaps: 22930-22958

-> Min/Max sizes since instance startup

-> Oper Types/Modes: INItializing, GROW, SHRink, STAtic/IMMediate, DEFerred

-> ordered by Component

| Component       | Begin Snap<br>Size (Mb) | Current<br>Size (Mb) | Min<br>Size (Mb) | Max<br>Size (Mb) | Oper<br>Count | Last Op<br>Typ/Mod |
|-----------------|-------------------------|----------------------|------------------|------------------|---------------|--------------------|
| ASM Buffer Cach | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 16K buf | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 2K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 32K buf | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 4K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT 8K buff | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| DEFAULT buffer  | 20,480.00               | 20,480.00            | 19,712.00        | 20,736.00        | 0             | GRO/DEF            |
| KEEP buffer cac | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| PGA Target      | 14,848.00               | 14,848.00            | 14,848.00        | 14,848.00        | 0             | STA/               |
| RECYCLE buffer  | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| SGA Target      | 28,160.00               | 28,160.00            | 28,160.00        | 28,160.00        | 0             | STA/               |
| Shared IO Pool  | .00                     | .00                  | .00              | .00              | 0             | STA/               |
| java pool       | 1,024.00                | 1,024.00             | 1,024.00         | 1,024.00         | 0             | SHR/DEF            |
| large pool      | 256.00                  | 256.00               | 256.00           | 256.00           | 0             | STA/               |
| shared pool     | 4,096.00                | 4,096.00             | 3,840.00         | 4,864.00         | 0             | SHR/DEF            |
| streams pool    | 2,048.00                | 2,048.00             | 2,048.00         | 2,048.00         | 0             | STA/               |

Texas Memory Systems, Inc.

The World's Fastest Storage®





V\$SGA\_RESIZE\_OPS -10g  
V\$MEMORY\_RESIZE\_OPS - 11g

- A dynamic performance view
- Use to get details of resize operations
- Every memory change since startup



# V\$SGA\_RESIZE\_OPS

Date: 02/18/08

Page: 1

Time: 01:04 PM

Component Resize Operations

SYSTEM

ault10g1 database

| COMPONENT            | Oper   | OPER_MODE | INITIAL_SIZE | TARGET_SIZE | FINAL_SIZE | STATUS   | START_TIME | END_TIME   |
|----------------------|--------|-----------|--------------|-------------|------------|----------|------------|------------|
| shared pool          | SHRINK | DEFERRED  | 318767104    | 301989888   | 301989888  | COMPLETE | 0217 15:23 | 0217 15:23 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1040187392   | 1056964608  | 1056964608 | COMPLETE | 0217 15:23 | 0217 15:23 |
| shared pool          | SHRINK | DEFERRED  | 301989888    | 285212672   | 285212672  | COMPLETE | 0217 15:24 | 0217 15:24 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1073741824   | 1090519040  | 1090519040 | COMPLETE | 0217 15:24 | 0217 15:24 |
| shared pool          | SHRINK | DEFERRED  | 285212672    | 268435456   | 268435456  | COMPLETE | 0217 15:24 | 0217 15:24 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1056964608   | 1073741824  | 1073741824 | COMPLETE | 0217 15:24 | 0217 15:24 |
| shared pool          | SHRINK | DEFERRED  | 268435456    | 251658240   | 251658240  | COMPLETE | 0217 15:25 | 0217 15:25 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1090519040   | 1107296256  | 1107296256 | COMPLETE | 0217 15:25 | 0217 15:25 |
| shared pool          | SHRINK | DEFERRED  | 251658240    | 234881024   | 234881024  | COMPLETE | 0217 15:28 | 0217 15:28 |
| DEFAULT buffer cache | GROW   | DEFERRED  | 1107296256   | 1124073472  | 1124073472 | COMPLETE | 0217 15:28 | 0217 15:28 |
| shared pool          | SHRINK | DEFERRED  | 234881024    | 218103808   | 218103808  | COMPLETE | 0217 15:28 | 0217 15:28 |

Texas Memory Systems, Inc.

The World's Fastest Storage®



# Parameters

init.ora Parameters

DB/Inst: Snaps: 84084-84108

| Parameter Name | Begin value | End value<br>(if different) |
|----------------|-------------|-----------------------------|
|----------------|-------------|-----------------------------|

|                           |             |  |
|---------------------------|-------------|--|
| _spin_count               | 2000        |  |
| cursor_sharing            | EXACT       |  |
| db_block_size             | 16384       |  |
| db_cache_size             | 46170898432 |  |
| db_keep_cache_size        | 5368709120  |  |
| db_name                   | cc1         |  |
| db_writer_processes       | 4           |  |
| java_pool_size            | 167772160   |  |
| large_pool_size           | 1073741824  |  |
| olap_page_pool_size       | 33554432    |  |
| open_cursors              | 900         |  |
| pga_aggregate_target      | 5368709120  |  |
| processes                 | 700         |  |
| sessions                  | 1000        |  |
| sga_max_size              | 68719476736 |  |
| sga_target                | 68719476736 |  |
| shared_pool_reserved_size | 262144000   |  |
| shared_pool_size          | 4294967296  |  |

Texas Memory Systems, Inc.

The World's Fastest Storage®



## Quick Word on RAC

- The TCP buffers are the biggest memory issue with RAC performance, size them big
- The global dictionary takes memory out of the shared pool, the more memory in the db caches, the more is needed in the shared pools to track it.



# Questions/Comments?



Texas Memory Systems, Inc.

The World's Fastest Storage®



Thank You!

Mike Ault

[Oracle@RamSan.com](mailto:Oracle@RamSan.com)

[www.ramsan.com](http://www.ramsan.com)

[www.statspackanalyzer.com](http://www.statspackanalyzer.com)

Texas Memory Systems, Inc.

— The World's Fastest Storage®