


■ ■ ■ ■ ■ Automatic Workload Repository



DOAG Jahreskonferenz 2004
10/11 November 2004, Mannheim

Christian Antognini
christian.antognini@trivadis.com

Trivadis AG
www.trivadis.com

trivadis
makes **IT** easier. ■ ■ ■

Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

Important Questions That Need Answers



- > The tuning of a whole system is much more complex than the tuning of a single component (e.g. a report or a user action)
 - ➔ Tune single components whenever possible

- > For each component the following questions need answers
 - ① Where is the time spent?
 - ② How is the time spent?
 - ③ How to reduce the time spent?

} Identification

} Tuning

Response Time Tuning



Response Time = Service Time + Wait Time

Cumulated Statistics – The Problem



- > Dynamic performance views like V\$SYSSTAT shows cumulative statistics since database startup
 - » E.g. number of logical I/Os

```
SQL> SELECT sum(value) lio
      2 FROM v$sysstat
      3 WHERE name IN ('db block gets','consistent gets');

          LIO
-----
57664003
```

- > Cumulated statistics are not very useful
 - ➔ We need a history
 - ➔ We need rates

Cumulated Statistics – The Solution



- > What is a rate?
 - » Take value v_1 at time t_1
 - » Take value v_2 at time t_2
 - » $\text{Rate} = (v_2 - v_1) / (t_2 - t_1)$

- > Statspack provides history and rates
 - » Installation required
 - » Overhead
 - » Limited snapshots
 - » Difficult trend analysis

- > Thanks to Automatic Workload Repository (AWR) in Oracle 10^g history and rates are automatically gathered and partially available in memory as well
 - » In AWR rates are named metrics

Automatic Workload Repository



- > The AWR stores performance statistics used for diagnostic and self-tuning purposes
 - » Time model
 - » Wait events
 - » Metrics
 - » Active session history
 - » OS statistics
 - » SQL statistics (e.g. logical reads per SQL statement)

- > Data is both in memory and stored in the database
 - » Recent performance statistics are available through V\$ views
 - » A history over several days is available through DBA_HIST views
 - » To reduce the required space, for some statistics, only alerts are stored

Automatic Workload Repository vs. Statspack



- > AWR is an enhanced Statspack
- > AWR is automatically installed and configured
- > AWR has new statistics thanks to 10_g features
- > AWR automatically takes snapshots
 - » Default 30 minutes
- > AWR automatically purge collected data
 - » Default 7 days
- > AWR has a GUI (Database/Grid Control)
 - » Trivadis has developed a GUI for Statspack: TVD\$STAT
- > Statspack still exists and has been enhanced as well

Snapshots Interval and Retention Time



- > They are shown in DBA_HIST_WR_CONTROL

```
SQL> SELECT snap_interval, retention
       2 FROM dba_hist_wr_control;
```

SNAP_INTERVAL	RETENTION
-----	-----
+00000 00:30:00.0	+00007 00:00:00.0

- > They can be modified through DBMS_WORKLOAD_REPOSITORY (in minutes)

```
SQL> BEGIN
       2 dbms_workload_repository.modify_snapshot_settings(
       3     interval => 60,
       4     retention => 10*24*60
       5 );
       6 END;
       7 /
```

Automatic Workload Repository Reports



- > They can be generated in text or HTML
 - » Database/Grid Control
 - » Package DBMS_WORKLOAD_REPOSITORY
 - » Scripts awrrpt.sql (current DBID and instance) or awrrpti.sql (DBID and instance are requested)
- > Very similar to Statspack reports
- > The script awrinfo.sql generates general information about AWR such as the size and data distribution

Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

Time Model Statistics



- > Two new views give high-level information about the used CPU and elapsed time for different database operations
 - » V\$SYS_TIME_MODEL
 - » V\$SESS_TIME_MODEL
 - » The amount of time is given in microseconds
- > The statistic “DB time” represents the total time spent in database calls; this doesn’t include the time spent by background processes
- > Since waits are part of the elapsed time, “DB time” can be much larger than the elapsed time since instance startup
- > Some of the information provided by these views is already available in V\$SYSTAT and V\$SESSTAT

Time Model Statistics – Components



background elapsed time

└─ background cpu time

DB time

└─ DB CPU

└─ connection management call elapsed time

└─ failed parse elapsed time

└─ failed parse (out of shared memory) elapsed time

└─ inbound PL/SQL rpc elapsed time

└─ Java execution elapsed time

└─ parse time elapsed

└─ hard parse elapsed time

└─ hard parse (sharing criteria) elapsed time

└─ hard parse (bind mismatch) elapsed time

└─ PL/SQL compilation elapsed time

└─ PL/SQL execution elapsed time

└─ sequence load elapsed time

└─ sql execute elapsed time

Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

What Is the Wait Interface?



- > Since version 7.0 the Oracle code has been instructed to report the time waited completing internal operations
- > Oracle calls these internal operations events or waits
- > The gathered information is available through V\$ views and trace facilities (e.g. event 10046)
- > The wait interface is the most important source of information used to analyze and/or discover performance problems

What Is the Wait Interface?



- > In the following example a TVD\$XTAT output shows that a statement “waited” about 17 seconds (difference between CPU and elapsed time)

Call	Misses	Count	CPU [s]	Elapsed [s]	PIO [b]	LIO[b]	Consistent [b]	Current [b]	Rows
Parse	1	1	0.010	0.014	0	0	0	0	0
Execute	0	1	0.120	17.699	1	63	7	56	17
Fetch	0	0	0.000	0.000	0	0	0	0	0
Total	1	2	0.130	17.713	1	63	7	56	17
Average (per row)	0	0	0.008	1.042	0	3	0	3	1

- > Where does this difference come from?

What Is the Wait Interface?



- > If the trace file has been generated with extended SQL trace information, e.g. through DBMS_MONITOR, the answer can be found in the TVD\$XTAT output as well

Component	Total Duration [s]	%	Number of Events	Duration per Event [s]
eng: TM - contention	17.357	82.382	6	2.893
SQL*Net message from client	2.921	13.862	1	2.921
recursive statements	0.490	2.326	n/a	n/a
db file sequential read	0.171	0.813	1	0.171
CPU	0.130	0.617	n/a	n/a
SQL*Net message to client	0.000	0.000	1	0.000
Total	21.069	100.000		

→ The statement waited about 17 seconds acquiring a row lock

- > This information comes from the wait interface

What Is New in the Wait Interface?



- > More events are available:
 - » 10.1.0.3 → 811
 - » 9.2.0.5 → 402
 - » 8.1.7.4 → 217
 - » 7.3.4.5 → 105
- > Waits have been grouped in 12 classes
(e.g. Application, Cluster, Concurrency, Network, User I/O)
- > More V\$ views contains wait interface information
(e.g. V\$SESSION contains the information that was previously available in V\$SESSION_WAIT)
- > New V\$ views

Wait Classes




- > In V\$EVENT_NAME new columns specify the wait class
- > V\$SESSION_WAIT_CLASS displays for each session the wait time (in centiseconds) grouped by class

```
SQL> SELECT wait_class, total_waits, time_waited
       2 FROM v$session_wait_class
       3 WHERE sid = 17;
```


WAIT_CLASS	TOTAL_WAITS	TIME_WAITED
-----	-----	-----
Idle	466	31640
User I/O	1613	317
System I/O	46	3
Application	2	1
Network	135	0

- > V\$SYSTEM_WAIT_CLASS and V\$SERVICE_WAIT_CLASS display wait class information for the system/services since startup

Wait Classes in Database/Grid Control



Drill down



Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

Which Metrics Exists?



- > Metrics are available for
 - » System
 - » Sessions
 - » Services
 - » Events
 - » Files
- > They are organized in groups
- > They are computed over an interval in the following units
 - » Absolute values
 - » Percentages
 - » Per second and per transaction
- > For each metric a warning and critical threshold can be configured

Which Metrics Exists?



```
SQL> SELECT g.group_id id, g.name, g.interval_size, g.max_interval,
2         count(*) cnt
3   FROM v$metricgroup g, v$metricname n
4  WHERE g.group_id = n.group_id
5  GROUP BY g.group_id, g.name, g.interval_size, g.max_interval;
```

ID	NAME	INTERVAL_SIZE	MAX_INTERVAL	CNT
0	Event Metrics	6000	1	3
1	Event Class Metrics	6000	60	4
2	System Metrics Long Duration	6000	60	119
3	System Metrics Short Duration	1500	12	40
4	Session Metrics Long Duration	6000	60	1
5	Session Metrics Short Duration	1500	1	8
6	Service Metrics	6000	60	2
7	File Metrics Long Duration	60000	6	6
9	Tablespace Metrics Long Duration	6000	0	1

Number of metrics for each "statistic",
e.g. each event has 3 metrics

Which Metrics Exists?



```
SQL> SELECT group_id id, metric_name, metric_unit
2   FROM v$metricname
3   ORDER BY group_id;
```


ID	METRIC_NAME	METRIC_UNIT
0	Total Wait Counts	Waits
0	Total Time Waited	CentiSeconds
0	Number of Sessions Waiting (Event)	Sessions
1	Total Wait Counts	Waits
1	Total Time Waited	CentiSeconds
1	Database Time Spent Waiting (%)	% (TimeWaited / DBTime)
1	Average Users Waiting Counts	Users
2	Host CPU Utilization (%)	% Busy/(Idle+Busy)
2	Database Time Per Sec	CentiSeconds Per Second
2	Txns Per Logon	Txns Per Logon
2	Executions Per Sec	Executes Per Second
...		
7	Average File Read Time (Files-Long)	CentiSeconds Per Read
9	Tablespace Space Usage	Tablespace Blocks

184 rows selected.

Current and History Rates – File Metrics Example



I/O – Cumulated value (V\$FILESTAT)



Views



> Information

- » V\$METRICNAME
- » V\$METRICGROUP
- » DBA_HIST_METRIC_NAME

> Global (contains all metrics)

- » V\$METRIC
- » V\$METRIC_HISTORY

> Service

- » V\$SERVICEMETRIC
- » V\$SERVICEMETRIC_HISTORY

> Session

- » V\$SESSMETRIC
- » DBA_HIST_SESSMETRIC_HISTORY^(a)

> System

- » V\$SYSMETRIC
- » V\$SYSMETRIC_HISTORY
- » V\$SYSMETRIC_SUMMARY
- » DBA_HIST_SYSMETRIC_HISTORY^(a)
- » DBA_HIST_SYSMETRIC_SUMMARY

> Events

- » V\$EVENTMETRIC
- » V\$WAITCLASSMETRIC
- » V\$WAITCLASSMETRIC_HISTORY

> Files

- » V\$FILEMETRIC
- » V\$FILEMETRIC_HISTORY
- » DBA_HIST_FILEMETRIC_HISTORY^(a)

Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

Active Session History



- > Active Session History (ASH) is a history of the information similar to that available in V\$SESSION and V\$SESSION_WAIT
 - » V\$SESSION_WAIT shows on which event a session is waiting
 - » V\$SESSION_WAIT_HISTORY shows the last 10 waits
- > A session is active if it uses CPU or waits for an event that doesn't belong to the wait class "Idle"
- > Active sessions are sampled every second and stored in a circular buffer in the SGA
 - » V\$ACTIVE_SESSION_HISTORY
 - » The retention is extremely variable
- > History available through DBA_HIST_ACTIVE_SESS_HISTORY
 - » Only a subset of V\$ACTIVE_SESSION_HISTORY is stored on disk

Session Wait History – Example



- > The following statement shows for what events waited the foreground sessions during the last hour

```
SQL> SELECT event, count(*), sum(time_waited) time_waited
  2  FROM v$active_session_history
  3  WHERE sample_time >= sysdate - to_dsinterval('0 1:0:0')
  4  AND session_type = 'FOREGROUND'
  5  GROUP BY event
  6  ORDER BY 3 DESC;
```

EVENT	COUNT(*)	TIME_WAITED
-----	-----	-----
enq: TX - row lock contention	525	517191348
enq: TM - contention	149	126135243
db file sequential read	61	3240729
latch: library cache	6	2262512
latch: enqueue hash chains	11	1622636
log file sync	34	1582910
latch free	1	549162
buffer busy waits	1	279213
...		

Automatic Workload Repository



> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

OS Statistics



- > The cumulated values are available in V\$OSSTAT
 - » Totals over all processors and average values are available
 - » The history is stored in DBA_HIST_OSSTAT

```
SQL> SELECT stat_name, value
       2 FROM v$osstat
       3 WHERE stat_name NOT LIKE 'AVG%';
```

STAT_NAME	VALUE
-----	-----
NUM_CPUS	1
IDLE_TICKS	35920275
BUSY_TICKS	2884344
USER_TICKS	950628
SYS_TICKS	1907011
NICE_TICKS	26705
RSRC_MGR_CPU_WAIT_TIME	0
IN_BYTES	4.7691E+10
OUT_BYTES	2.3344E+11

Performance Impact



- > Of course the gathering and processing of a large amount of performance statistics has a performance impact on the database
- > Usually this overhead can be neglected
 - » Oracle says that during a TCP benchmark the overhead was 3%
- > To disable AWR the INIT.ORA parameter STATISTICS_LEVEL can be set to BASIC

Licensing Information



- > To use the following features/components the Diagnostic Pack must be licensed
 - » Package DBMS_WORKLOAD_REPOSITORY
 - » View V\$ACTIVE_SESSION_HISTORY
 - » All views with the prefix DBA_HIST
 - » Report awrrpt.sql and awrrpti.sql
- > This means that per default when a database is created it contains features/components that cannot be used!
- > For more information refer to the document *Oracle Database Licensing Information 10g Release 1*
 - » The document also describes on which links of the Database/Grid Control you can click...



Automatic Workload Repository




> Agenda

Be in
the know.

- > Introduction
- > Time Model
- > Wait Interface
- > Metrics
- > Active Session History
- > Miscellaneous
- > Conclusion

Automatic Workload Repository: Core Statements...




- + A lot of very interesting performance statistics with history and rates
- ! Such a large amount of data increase complexity
 - » Which view contains the information I'm looking for?
- Questionable licensing strategies...

> from Trivadis

At the core it's
all about data.

■ ■ ■ ■ ■ Thank you for your attention



Questions?

trivadis
makes IT easier. ■ ■ ■