

Oracle SQL Plan Execution: How It Really Works and How to Troubleshoot It


Tanel Põder

<http://blog.tanelpoder.com>

<http://www.e2sn.com>

Introduction

About me:

- Occupation: DBA, researcher, consultant
- Expertise: Oracle internals geek,
End-to-end performance & scalability,
Oracle troubleshooting,
Oracle capacity planning
- Oracle experience: 13+ years as DBA
- Certification: OCM (2002) OCP (1999)
- Professional affiliations: OakTable Network 
Oracle ACE Director
- Blog: <http://blog.tanelpoder.com>
- Company: <http://www.e2sn.com>



What is an Oracle workload about?

A bunch of sessions

- You need to have a session for doing anything in an Oracle database
- Every session has entries in V\$SESSION, V\$SESSTAT, etc

Executing SQL (or PL/SQL) cursors

- Every SQL has a SQL_ID (or hash_value which essentially is the same thing)
- Even PL/SQL calls use a cursor for invoking the stored procedure
 - Since Oracle 10.2.0.3 Oracle reports the PLSQL_OBJECT/SUBPROGRAM ID in V\$SESSION

Running on CPU - or not running on CPU

- Running on CPU
 - Wanting to run on CPU (but OS doesn't allow it)
- Not running on CPU
 - Sleeping, waiting for system call to complete

What is an execution plan?

For Oracle server:

- *Parsed*, optimized and compiled SQL code kept inside library cache

For DBAs and developers:

- Text or graphical representation of SQL execution flow

Often known as *explain plan*

- To be correct in terms, explain plan is just a tool, command in Oracle
- Explain plan outputs textual representation of execution plan into plan table
- DBAs/developers report human readable output from plan table

Viewing execution plans

DBMS_XPLAN

- *Explain plan for select*
select * from table(dbms_xplan.display())
- select * from table(dbms_xplan.display_cursor(null,null, 'allstats'))
- select * from table(dbms_xplan.display_cursor(<sqlid>, <child>, 'advanced'))
- select * from table(dbms_xplan.display_awr(<sqlid>))
- @x.sql

V\$SQL_PLAN (and V\$SQL_PLAN_STATISTICS[_ALL])

- @xms
- @xmsh <hash_value> <child#>

V\$SQL_MONITOR / V\$SQL_PLAN_MONITOR (11g+)

- DBMS_SQLTUNE.REPORT_SQL_MONITOR
- @xp <SID> or @xph <SID>

Requires Oracle
diagnostics and tuning
pack license

event 10132 level 1

- Dumps execution plan to trace file every hard parse

Parse stages

Syntactic check

- Syntax, keywords, sanity

Semantic check

- Whether objects referenced exist, are accessible (by permissions) and are usable

View merging

- Queries are written to reference base tables
- Can merge both stored views and inline views

Query transformation

- Transitivity, etc (example: if $a=1$ and $a=b$ then $b=1$)

Optimization

Query execution plan (QEP) generation

Loading SQL and execution plan in library cache

soft parse

hard parse

View merging

Optimizer merges subqueries, inline and stored views and runs queries directly on base tables

- Not always possible though due semantic reasons

```
SQL> create or replace view empview
  2  as
  3  select e.empno, e.ename, d.dname
  4  from emp e, dept d
  5  where e.deptno = d.deptno;
```

```
SQL> select * from empview
  2  where ename = 'KING';
```

Can be controlled using:
Parameter: `_complex_view_merging`
 `_simple_view_merging`
Hints: MERGE, NO_MERGE

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		7	210	5 (20)
* 1	HASH JOIN		7	210	5 (20)
2	TABLE ACCESS FULL	DEPT	4	52	2 (0)
* 3	TABLE ACCESS BY INDEX ROWID	EMP	7	119	2 (0)
* 4	INDEX RANGE SCAN	EMP_ENAME	8		1 (0)

Subquery unnesting

Subqueries can be unnested, converted to anti- and semijoins

```
SQL> select * from employees e
      2 where exists (
      3     select ename from bonus b
      4     where e.ename = b.ename
      5 );
```

Can be controlled using:

Parameter: `_unnest_subqueries`

Hints: `UNNEST, NO_UNNEST`

Id	Operation	Name	Rows	Bytes	Cost
0	SELECT STATEMENT		1	37	5
1	NESTED LOOPS		1	37	5
2	NESTED LOOPS		1	24	4
3	SORT UNIQUE		1	7	2
4	TABLE ACCESS FULL	BONUS	1	7	2
* 5	TABLE ACCESS BY INDEX ROWID	EMP	1	17	1
* 6	INDEX RANGE SCAN	EMP_ENAME	37		1
7	TABLE ACCESS BY INDEX ROWID	DEPT	1	13	1
* 8	INDEX UNIQUE SCAN	PK_DEPT	1		0

Predicate Information (identified by operation id):

- 5 - filter("E"."DEPTNO" IS NOT NULL)
- 6 - access("E"."ENAME"="B"."ENAME")
- 8 - access("E"."DEPTNO"="D"."DEPTNO")

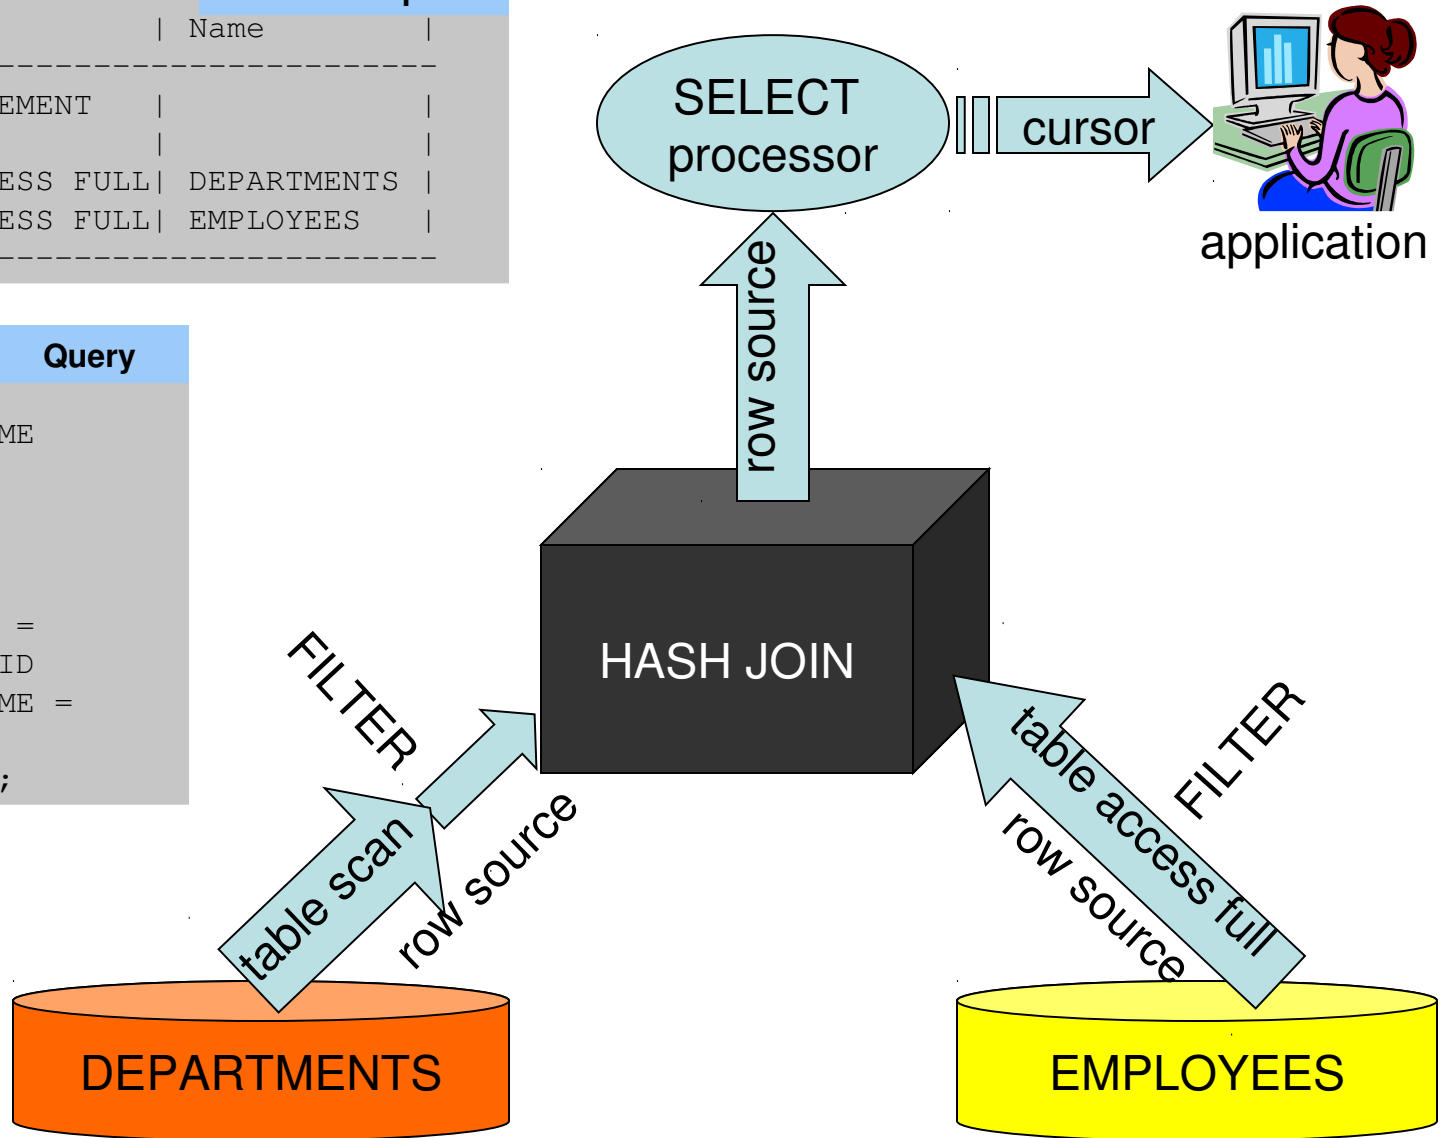
SQL execution basics

Execution plan

Id	Operation	Name
0	SELECT STATEMENT	
* 1	HASH JOIN	
* 2	TABLE ACCESS FULL	DEPARTMENTS
* 3	TABLE ACCESS FULL	EMPLOYEES

Query

```
SELECT
  E.LAST_NAME,
  D.DEPARTMENT_NAME
FROM
  EMPLOYEES E,
  DEPARTMENTS D
WHERE
  E.DEPARTMENT_ID =
    D.DEPARTMENT_ID
AND D.DEPARTMENT_NAME =
  'Sales'
AND E.SALARY > 2000;
```

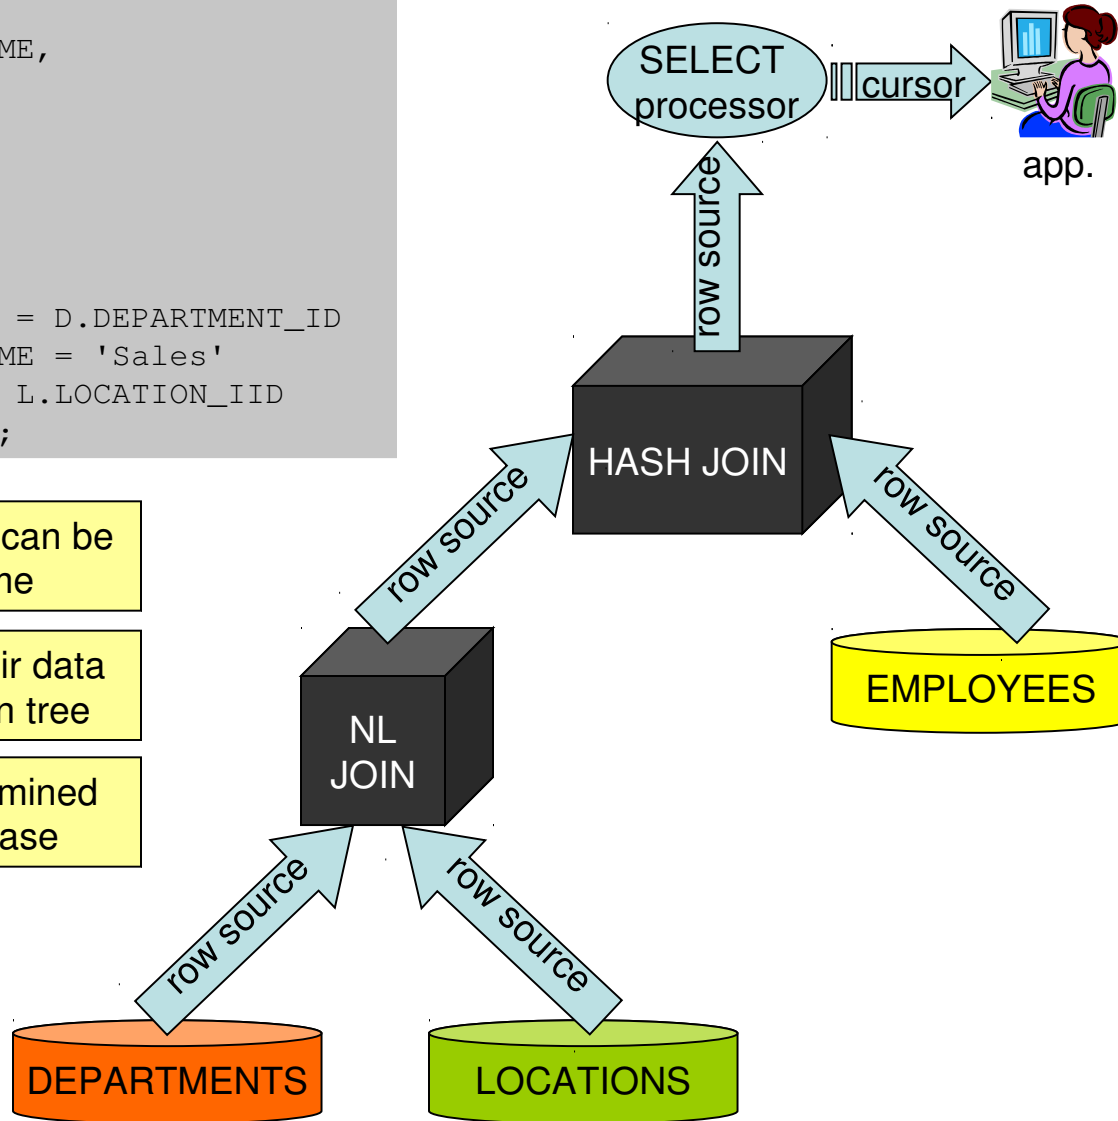


SQL execution basics - multitable joins

```
SELECT
  E.LAST_NAME,
  D.DEPARTMENT_NAME,
  L.CITY
FROM
  EMPLOYEES E,
  DEPARTMENTS D,
  LOCATIONS L
WHERE
  E.DEPARTMENT_ID = D.DEPARTMENT_ID
AND D.DEPARTMENT_NAME = 'Sales'
AND D.LOCATION_ID = L.LOCATION_ID
AND E.SALARY > 2000;
```

Multiple joins

- Only two row sources can be joined together at a time
- Row sources pass their data "up" the execution plan tree
- The join order is determined during optimization phase



SQL execution terminology

ACCESS PATH

- A means to access physical data in database storage
- From tables, indexes, external tables, database links

ROW SOURCE

- A virtual stream of rows
- Can come through access paths from tables, indexes
- Or from other child row sources

FILTER *PREDICATE*

- A property of row source - can discard rows based on defined conditions - *filter predicates*

JOIN

- Filters and merges rows based on matching rows from child row sources. Matching is defined by *join predicates*
- Any join operator can join only two inputs

First rule for reading an execution plan

Parent operations get input only from their children

Id Operation		Execution plan structure
0	SELECT STATEMENT	
* 1	FILTER	
2	NESTED LOOPS OUTER	
* 3	HASH JOIN OUTER	
4	NESTED LOOPS OUTER	
5	NESTED LOOPS OUTER	
* 6	HASH JOIN	
7	TABLE ACCESS FULL	USER\$
8	NESTED LOOPS	
* 9	HASH JOIN	
10	MERGE JOIN CARTESIAN	
* 11	HASH JOIN	
* 12	FIXED TABLE FULL	X\$KSPPI
13	FIXED TABLE FULL	X\$KSPPCV
14	BUFFER SORT	
15	TABLE ACCESS FULL	TS\$
* 16	TABLE ACCESS FULL	TAB\$
* 17	TABLE ACCESS BY INDEX ROWID	OBJ\$
* 18	INDEX UNIQUE SCAN	I_OBJ1
19	TABLE ACCESS BY INDEX ROWID	OBJ\$
* 20	INDEX UNIQUE SCAN	I_OBJ1
21	TABLE ACCESS BY INDEX ROWID	OBJ\$
* 22	INDEX UNIQUE SCAN	I_OBJ1
23	TABLE ACCESS FULL	USER\$
24	TABLE ACCESS CLUSTER	SEG\$
* 25	INDEX UNIQUE SCAN	I_FILE#_BLOCK#
26	NESTED LOOPS	
* 27	INDEX RANGE SCAN	I_OBJAUTH1
* 28	FIXED TABLE FULL	X\$KZSRO
* 29	FIXED TABLE FULL	X\$KZSPR

Second rule for reading an execution plan

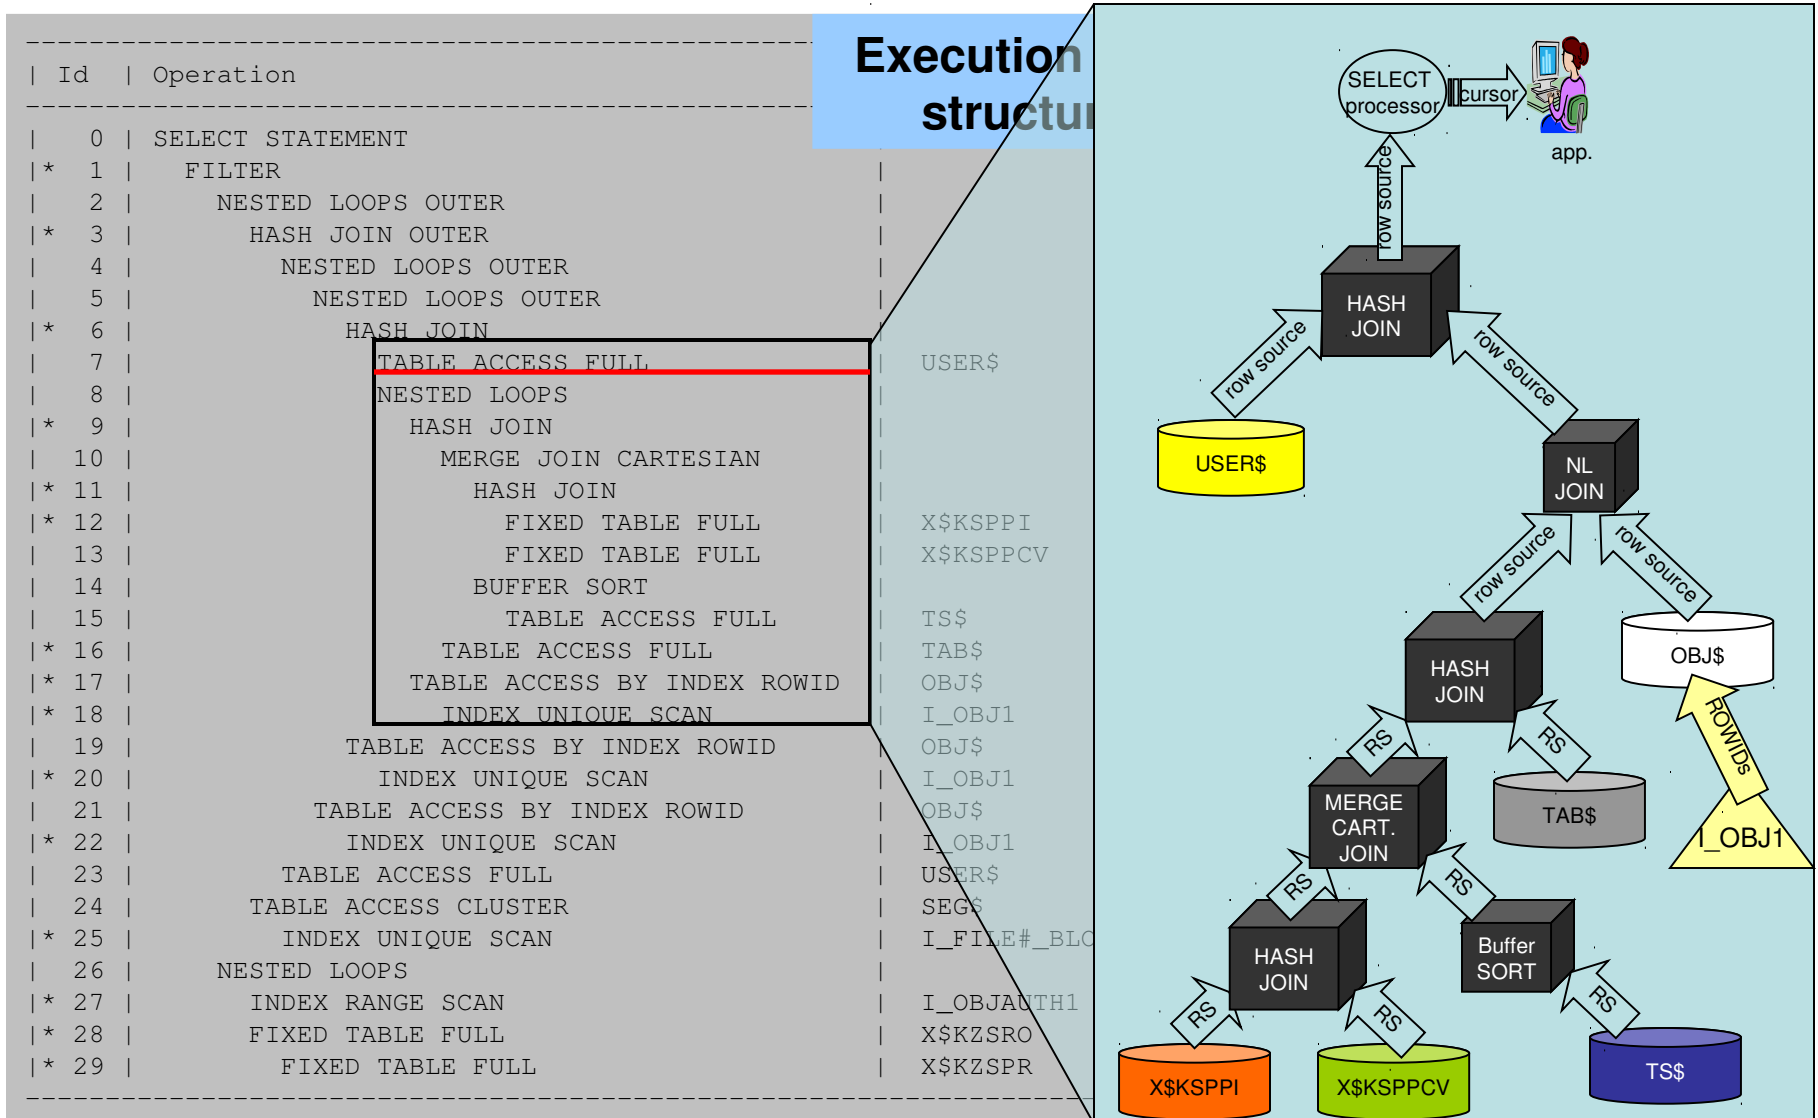
Data access starts from the first line without children

Execution plan structure	
Id	Operation
0	SELECT STATEMENT
* 1	FILTER
2	NESTED LOOPS OUTER
* 3	HASH JOIN OUTER
4	NESTED LOOPS OUTER
5	NESTED LOOPS OUTER
* 6	HASH JOIN
7	TABLE ACCESS FULL
8	NESTED LOOPS
* 9	HASH JOIN
10	MERGE JOIN CARTESIAN
* 11	HASH JOIN
* 12	FIXED TABLE FULL
13	FIXED TABLE FULL
14	BUFFER SORT
15	TABLE ACCESS FULL
* 16	TABLE ACCESS FULL
* 17	TABLE ACCESS BY INDEX ROWID
* 18	INDEX UNIQUE SCAN
19	TABLE ACCESS BY INDEX ROWID
* 20	INDEX UNIQUE SCAN
21	TABLE ACCESS BY INDEX ROWID
* 22	INDEX UNIQUE SCAN
23	TABLE ACCESS FULL
24	TABLE ACCESS CLUSTER
* 25	INDEX UNIQUE SCAN
26	NESTED LOOPS
* 27	INDEX RANGE SCAN
* 28	FIXED TABLE FULL
* 29	FIXED TABLE FULL

First operation with no children (leaf operation) **accesses** data

Cascading rowsources

Rows "flow" upwards to parent rowsources in cascading fashion



SQL execution plan recap

Execution plan lines are just Oracle kernel functions!

- In other words, each row source is a function

Data can only be accessed using *access path functions*

- Only access paths can access physical data
- Access paths process physical data, return *row sources*

Data processing starts from first line without children

- In other words the first leaf access path in execution plan

Row sources feed data to their parents

- Can be non-cascading, semi-cascading or cascading

A join operation can input only two row sources

- However, it is possible to combine result of more than 2 row sources for some operations (not for joins though)
- Index combine, bitmap merging, filter, union all, for example

SQL Plan profiling

SQL execution plan line level profiling available since Oracle 9.2

- Stats externalized in V\$SQL_PLAN_STATISTICS[_ALL]

Statistics gathering is enabled by setting parameter:

- statistics_level=all
- ...or _rowsource_execution_statistics=true
- or via hint: /*+ gather_plan_statistics */ (Oracle 10.2+)
- Don't enable this at instance level as it can kill your performance

```
$ pstack 1780 | ./os_explain
kpoal8
  SELECT FETCH:
    QUERY EXECUTION STATISTICS: Fetch
    GROUP BY SORT: Fetch
    QUERY EXECUTION STATISTICS: Fetch
    NESTED LOOP JOIN: Fetch
    QUERY EXECUTION STATISTICS: Fetch
    SORT: Fetch
    sorgetqbf
```

- Parameter introduced for reducing profiling overhead via reducing gettimeofday() syscalls
 - _rowsource_statistics_sampfreq = 128

Reading DBMS_XPLAN execution plan profile

```
SQL> select * from table(dbms_xplan.display_cursor(null,null,'ALLSTATS LAST'));
PLAN_TABLE_OUTPUT
```

```
-----
SQL_ID  56bs32ukywdsq, child number 0
-----
```

```
select count(*) from dba_tables
```

```
Plan hash value: 736297560
```

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time
1	SORT AGGREGATE		1	1	1	00:00:00.38
* 2	HASH JOIN RIGHT OUTER		1	1690	1688	00:00:00.37
3	TABLE ACCESS FULL	USERS\$	1	68	68	00:00:00.01
* 4	HASH JOIN OUTER		1	1690	1688	00:00:00.37
* 5	HASH JOIN					
6	TABLE ACCESS FU					
* 7	HASH JOIN					
8	NESTED LOOPS O					
* 9	HASH JOIN RIG					
10	TABLE ACCESS					
* 11	HASH JOIN					
12	MERGE JOIN					
* 13	HASH JOIN					
* 14	FIXED TAB					
15	FIXED TAB					
16	BUFFER SOR					
17	TABLE ACC					
* 18	TABLE ACCES					
* 19	INDEX UNIQUE					
* 20	TABLE ACCESS F					
21	TABLE ACCESS FULL					

Starts number of times the rowsource was initialized

E-rows CBO number estimated rows coming from rowsource

A-rows actual *measured* number of rows during last execution

A-time actual *measured (and extrapolated)* time spent inside a rowsource function or under its children (cumulative)

Buffer number of buffer gets done within rowsource during last execution

Reading XMS/XMSH execution plan profile

SQL> @xms

SQL hash value: 2783852310 Cursor address: 00000003DCA9EF28 | Statement first

Ch	Pr	Op	Object	ms spent	Estimated	Real #rows	Op. ite-
ld	ed	ID Operation	Name	in op.	output rows	returned	rations
0		0 SELECT STATEMENT					
		1 SORT AGGREGATE					
A		2 HASH JOIN					
		3 TABLE ACCESS					
A		4 HASH JOIN					
A		5 HASH JOIN					
		6 TABLE ACCESS					
A		7 HASH JOIN					
		8 NESTED LOOP					
A		9 HASH JOIN					
		10 TABLE ACCESS					
A		11 HASH JOIN					
		12 MERGE JOIN					
A		13 HASH JOIN					
F		14 TABLE ACCESS					
		15 TABLE ACCESS					
		16 TABLE ACCESS					
F		18 TABLE ACCESS					
A		19 INDEX					
F		20 TABLE ACCESS					
		21 TABLE ACCESS					
Ch	Op	Predicate In					
ld	ID						
0	2	- access ("CX")					
	4	- access ("T")					
	5	- access ("O")					
	7	- access ("O")					
	9	- access ("T")					

ms spent in op.

milliseconds spent in rowsource function (cumulative)

Estimated rows

CBO rowcount estimate

Real # rows

Real *measured* rowcount from rowsource

Op. iterations

Number of times the rowsource was initialized

Logical reads

Consistent buffer gets

Logical writes

Current mode buffer gets (Note that some CUR gets may not always be due writing...)

Physical reads

Physical reads done by the rowsource function

Physical writes

Physical writes done by the rowsource function

Optimizer cost

Least significant thing for measuring the *real execution efficiency* of a statement

Real time SQL execution monitoring

Oracle 11g new feature

- Uses V\$SQL_MONITOR and V\$SQL_PLAN_MONITOR
- Always enabled for parallel execution queries
- Kicks in for serial queries after they've waited total 5 seconds for IO or have used CPU
 - *_sqlmon_threshold = 5*
- You can also use MONITOR and NO_MONITOR hints for controlling the monitoring

~~Get execution statistics of last query executed in session:~~

SELECT

```
DBMS_SQLTUNE.REPORT_SQL_MONITOR (  
    session_id=>sys_context('userenv','sid'),  
    report_level=>'ALL',
```

Execution Profile (dbms_sqltune.report_sql_monitor)

SQL> @xp 128

REPORT

SQL Monitoring Report

SQL Text

```
select /*+ ordered use_nl(b) full(a) full(b) */ count(*) from sys.obj$ a, sys.obj$ b where a.name = b.name and r
```

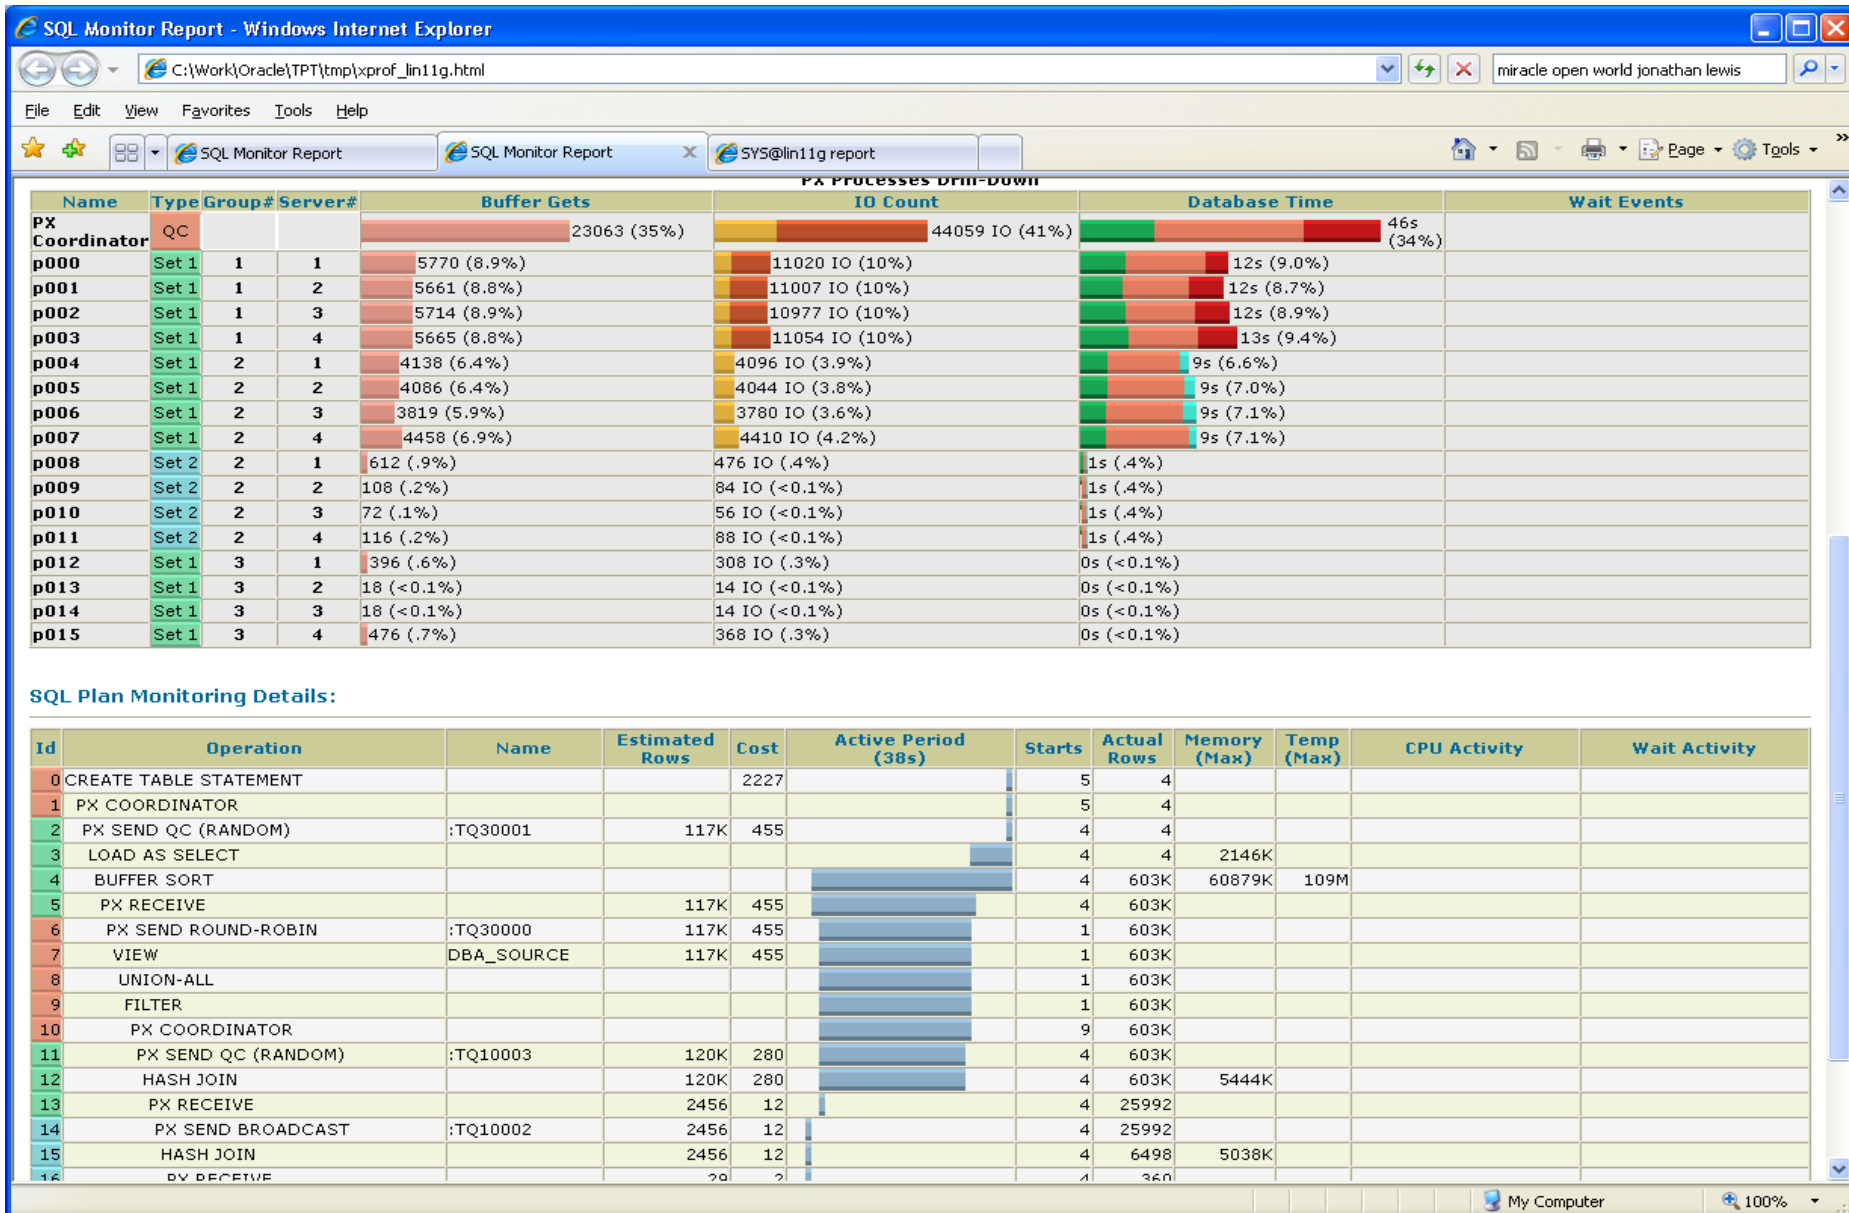
Global Information

Status : **EXECUTING**
Instance ID : 1
Session ID : 128
SQL ID : lvm188y2gv75n
SQL Execution ID : 16777217
Plan Hash Value : 2119813036
Execution Started : 08/14/2008 18:12:52
First Refresh Time : 08/14/2008 18:13:00
Last Refresh Time : 08/14/2008 18:13:20

SQL Plan Monitoring Details

Id	Operation	Name	Rows (Estim)	Cost	Time Active(s)	Start Active	Starts	Rows (Actual)	Activity (percent)
0	SELECT STATEMENT			16502K			1		
1	SORT AGGREGATE		1				1		
-> 2	COUNT STOPKEY				21	+8	1	3006	
-> 3	NESTED LOOPS		116K	16502K	21	+8	1	3006	
-> 4	TABLE ACCESS FULL	OBJ\$	69996	238	21	+8	1	2925	
-> 5	TABLE ACCESS FULL	OBJ\$	2	236	28	+1	2926	3006	100.00

Execution Profile HTML output (type=>'HTML')



Simple full table scan

Full table scan scans all the rows in the table

- All table blocks are scanned up to the HWM
- Even if all rows have been deleted from table
- Oracle uses multiblock reads where it can
- Most efficient way when querying majority of rows
 - And majority of columns

```
SQL> select * from emp;
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
Plan hash value: 4080710170
```

```
-----  
| Id  | Operation                | Name | Rows  | Bytes | Cost  (%CPU) | Time      |  
-----  
|   0 | SELECT STATEMENT         |      |    14 |    518 |    3   (0) | 00:00:01 |  
|   1 |   TABLE ACCESS FULL EMP | EMP  |    14 |    518 |    3   (0) | 00:00:01 |  
-----
```

Full table scan with a filter predicate

Filter operation throws away non-matching rows

- By definition, not the most efficient operation
- Filter conditions can be seen in predicate section

```
SQL> select * from emp where ename = 'KING';
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
Plan hash value: 4080710170
```

```
-----  
| Id  | Operation                | Name | Rows  | Bytes | Cost (%CPU)| Time     |  
-----  
|  0  | SELECT STATEMENT         |      |    1  |    37 |    3   (0)| 00:00:01 |  
|*   1  | TABLE ACCESS FULL      | EMP  |    1  |    37 |    3   (0)| 00:00:01 |  
-----
```

```
Predicate Information (identified by operation id):
```

```
-----  
1 - filter("ENAME"='KING')
```

Simple B*tree index+table access

Index tree is walked from root to leaf

- Key values and ROWIDs are gotten from index
- Table rows are gotten using ROWIDs
- Access operator fetches only matching rows
 - As opposed to *filter* which filters through the whole child rowsource

```
SQL> select * from emp where empno = 10;
```

Id	Operation	Name	Rows	Bytes	Cost	(%CPU)

0	SELECT STATEMENT		1	37	1	(0)
1	TABLE ACCESS BY INDEX ROWID	EMP	1	37	1	(0)
* 2	INDEX UNIQUE SCAN	PK_EMP	1		0	(0)

Predicate Information (identified by operation id):

2 - access ("EMPNO"=10)

Predicate attributes

Predicate = access

- A means to avoid processing (some) unneeded data at all

Predicate = filter

- Everything from child row source is processed / filtered
- The non-matching rows are *thrown away*

```
SQL> select * from emp
      2  where empno > 7000
      3  and ename like 'KING%';
```

Id	Operation	Name	Rows	Bytes	Cost	(%CPU)

0	SELECT STATEMENT		1	27	3	(0)
* 1	TABLE ACCESS BY INDEX ROWID	EMP	1	27	3	(0)
* 2	INDEX RANGE SCAN	PK_EMP	9		2	(0)

Predicate Information (identified by operation id):

PLAN_TABLE_OUTPUT

```
-----
1 - filter("ENAME" LIKE 'KING%')
2 - access("EMPNO">7000)
```

Index fast full scan

Doesn't necessarily return keys in order

- The whole index segment is just scanned as Oracle finds its blocks on disk (in contrast to tree walking)
- Multiblock reads are used
- As indexes don't usually contain all columns that tables do, FFS is more efficient if all used columns are in index
- Used mainly for aggregate functions, min/avg/sum,etc
- Optimizer must know that all table rows are represented in index! (null values and count example)

```
SQL> select min(empno), max(empno) from emp;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		1	5	25 (0)
1	SORT AGGREGATE		1	5	
2	INDEX FAST FULL SCAN	PK_EMP	54121	264K	25 (0)

Nested Loop Join

Nested loop join

- Read data from outer row source (upper one)
- *Probe* for a match in inner row source for each outer row

```
SQL> select d.dname, d.loc, e.empno, e.ename
2   from emp e, dept d
3   where e.deptno = d.deptno
4   and d.dname = 'SALES'
5   and e.ename like 'K%';
```

Id	Operation	Name	Rows	Bytes	Cost
0	SELECT STATEMENT		1	37	4
1	NESTED LOOPS		1	37	4
* 2	TABLE ACCESS FULL	EMP	1	17	3
* 3	TABLE ACCESS BY INDEX ROWID	DEPT	1	20	1
* 4	INDEX UNIQUE SCAN	PK_DEPT	1		

Predicate Information (identified by operation id):

- 2 - filter("E"."DEPTNO" IS NOT NULL AND "E"."ENAME" LIKE 'K%')
- 3 - filter("D"."DNAME"='SALES')
- 4 - access("E"."DEPTNO"="D"."DEPTNO")

Hash Join

Only for equijoins/non-equijoins (outer joins in 10g)

- Builds an array with hashed key values from smaller row source
- Scans the bigger row source, builds and compares hashed key values on the fly, returns matching ones

```
SQL> select d.dname, d.loc, e.empno, e.ename
2   from emp e, dept d
3   where e.deptno = d.deptno
4   and d.dname = 'SALES'
5   and e.ename between 'A%' and 'M%';
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	
0	SELECT STATEMENT		1	37	9 (12)	
* 1	HASH JOIN		1	37	9 (12)	
* 2	TABLE ACCESS FULL	DEPT	1	20	2 (0)	
* 3	TABLE ACCESS FULL	EMP	4	68	6 (0)	

Predicate Information (identified by operation id):

- 1 - access ("E"."DEPTNO"="D"."DEPTNO")**
- 2 - filter("D"."DNAME"='SALES')
- 3 - filter("E"."DEPTNO" IS NOT NULL AND "E"."ENAME"<='M%'
AND "E"."ENAME">='A%')

Sort-Merge Join

Requires both row sources to be sorted

- Either by a sort operation
- Or sorted by access path (index range and full scan)

Cannot return any rows before both row sources are sorted (non-cascading)

NL and Hash join should be normally preferred

```
SQL> select /*+ USE_MERGE(d,e) */ d.dname, d.loc, e.empno, e.ename
  2  from emp e, dept d
  3  where e.deptno = d.deptno
  4  and d.dname = 'SALES'
  5  and e.ename between 'A%' and 'X%'
  6  order by e.deptno;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		1245	46065	64 (10)
1	MERGE JOIN		1245	46065	64 (10)
* 2	TABLE ACCESS BY INDEX ROWID	DEPT	1	20	2 (0)
3	INDEX FULL SCAN	PK_DEPT	4		1 (0)
* 4	SORT JOIN		3735	63495	62 (10)
* 5	TABLE ACCESS FULL	EMP	3735	63495	61 (9)

Conclusion

Identifying problem SQL in the database

- Measure, don't guess!
- As easy as just querying V\$SESSION
 - Remember, a database workload is just a bunch of sessions, running SQL, waiting or working
- @a.sql
- @snapper_v3
- Perfsheet

Identifying the problem inside a SQL

- Measure, don't guess!
- Don't use just *explain plan*
- As it may show a wrong plan it doesn't show the real execution statistics
 - *Run* the statement and gather actual *execution statistics*
- Report with DBMS_XPLAN.DISPLAY_CURSOR or @xmsh.sql
- This allows profiling of where most of the response time has been spent
- Also compare real row counts vs estimated row counts

Advanced Oracle SQL Tuning

- *10-12. May, Singapore (3 days)*
- How to get systematic about Oracle SQL tuning

Parallel Execution and Partitioning for Performance

- *13. May, Singapore*
- How to get the best performance out of partitioning and parallel execution

Advanced Oracle Troubleshooting

- July 2010...
- 3 days of *intensive* database troubleshooting
- How to troubleshoot hangs, crashes, deadlocks, latch, lock contention, bugs and bad performance

<http://tech.e2sn.com/oracle-training-seminars>

Questions?

Tanel Põder

<http://blog.tanelpoder.com>

<http://www.e2sn.com>